TÉCNICO LISBOA

Flight Time and Cost Minimization in Complex Routes

Rafael Alexandre da Silva Marques

Thesis to obtain the Master of Science Degree in

Masters in Aerospace Engineering

Supervisors: Prof. Nuno Roma
Prof. Prof. Luis Russo

Examination Committee

Chairperson: Prof. Name of the Chairperson
Supervisor: Prof. Nuno Roma
Members of the Committee: Prof. Name of First Committee Member
Dr. Name of Second Committee Member
Eng. Name of Third Committee Member

October 2018

Acknowledgments

I would like to thank my parents for their friendship, encouragement and caring over all these years, for always being there for me through thick and thin and without whom this project would not be possible. I would also like to thank my grandparents, aunts, uncles and cousins for their understanding and support throughout all these years.

Quisque facilisis erat a dui. Nam malesuada ornare dolor. Cras gravida, diam sit amet rhoncus ornare, erat elit consectetuer erat, id egestas pede nibh eget odio. Proin tincidunt, velit vel porta elementum, magna diam molestie sapien, non aliquet massa pede eu diam. Aliquam iaculis.

Fusce et ipsum et nulla tristique facilisis. Donec eget sem sit amet ligula viverra gravida. Etiam vehicula urna vel turpis. Suspendisse sagittis ante a urna. Morbi a est quis orci consequat rutrum. Nullam egestas feugiat felis. Integer adipiscing semper ligula. Nunc molestie, nisl sit amet cursus convallis, sapien lectus pretium metus, vitae pretium enim wisi id lectus.

Donec vestibulum. Etiam vel nibh. Nulla facilisi. Mauris pharetra. Donec augue. Fusce ultrices, neque id dignissim ultrices, tellus mauris dictum elit, vel lacinia enim metus eu nunc.

I would also like to acknowledge my dissertation supervisors Prof. Some Name and Prof. Some Other Name for their insight, support and sharing of knowledge that has made this Thesis possible.

Last but not least, to all my friends and colleagues that helped me grow as a person and were always there for me during the good and bad times in my life. Thank you.

To each and every one of you – Thank you.

Abstract

Nulla facilisi. In vel sem. Morbi id urna in diam dignissim feugiat. Proin molestie tortor eu velit. Aliquam erat volutpat. Nullam ultrices, diam tempus vulputate egestas, eros pede varius leo, sed imperdiet lectus est ornare odio. Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Proin consectetuer velit in dui. Phasellus wisi purus, interdum vitae, rutrum accumsan, viverra in, velit. Sed enim risus, congue non, tristique in, commodo eu, metus. Aenean tortor mi, imperdiet id, gravida eu, posuere eu, felis. Mauris sollicitudin, turpis in hendrerit sodales, lectus ipsum pellentesque ligula, sit amet scelerisque urna nibh ut arcu. Aliquam in lacus. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Nulla placerat aliquam wisi. Mauris viverra odio. Quisque fermentum pulvinar odio. Proin posuere est vitae ligula. Etiam euismod. Cras a eros.

Keywords

Maecenas tempus dictum libero; Donec non tortor in arcu mollis feugiat;Cras rutrum pulvinar tellus.

Resumo

Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Vestibulum tortor quam, feugiat vitae, ultricies eget, tempor sit amet, ante. Donec eu libero sit amet quam egestas semper. Aenean ultricies mi vitae est. Mauris placerat eleifend leo. Quisque sit amet est et sapien ullamcorper pharetra. Vestibulum erat wisi, condimentum sed, commodo vitae, ornare sit amet, wisi. Aenean fermentum, elit eget tincidunt condimentum, eros ipsum rutrum orci, sagittis tempus lacus enim ac dui. Donec non enim in turpis pulvinar facilisis. Ut felis. Aliquam aliquet, est a ullamcorper condimentum, tellus nulla fringilla elit, a iaculis nulla turpis sed wisi. Fusce volutpat. Etiam sodales ante id nunc. Proin ornare dignissim lacus. Nunc porttitor nunc a sem. Sed sollicitudin velit eu magna. Aliquam erat volutpat. Vivamus ornare est non wisi. Proin vel quam. Vivamus egestas. Nunc tempor diam vehicula mauris. Nullam sapien eros, facilisis vel, eleifend non, auctor dapibus, pede.

Palavras Chave

Colaborativo; Codificaçãoo; Conteúdo Multimédia; Comunicação;

Contents

List of Figures

x

List of Tables

List of Algorithms

Listings

Acronyms

1

Introduction

Contents

1.1 Motivation

Online Traveling Agencies (OTA) are online applications which sell traveling goods, as, for example, commercial flight tickets. Although consumers retain the option to buy flights directly from airline companies, the majority chooses to use OTA. The main reason for this is that these agencies aggregate flight data from multiple airlines, instead of being limited to a single one, which ultimately increases the options of the consumer. Furthermore, many OTA work as meta-search engines, searching over a variety of websites in order to find the best flights which satisfy the consumer requirements. While OTA usually provide very complete search functionalities for simple flights, the majority fails to offer the same search options for a trip composed of multiple cities.

Consider a trip starting and ending at any given city, which must visit every city specified in a particular list of cities. If there are no constraints associated with the order in which these cities must be visited, this problem is known, in the scientific community, as the Traveling Salesman Problem (TSP). This problem is considered very difficult to solve, since the total number of possible solutions increases in an exponential way, as the number of cities increases.

If commercial flights are the means of transportation between every pair of cities, this problem can no longer be considered the *classic* Traveling Salesman, but rather its *time-dependent* counterpart. This is because some of the major flight characteristics, as its price and duration, cannot be considered constant over time. Rather, they are dependent on the particular flight selected, and the characteristics of that flight may follow no apparent logical rule, at least from the consumers perspective.

Consequently, finding the most efficient set of flights, from the consumer point of view, is a very repetitive and time-consuming task. Faced with this problem, only the most persistent consumer will be able to find the best solution for the problem, and this can only occur for a very small list of cities. As the number of cities increases, even the most persistent consumer cannot verify all possible solutions. This means that, ultimately, the final consumer will pay more than necessary for the requested service.

This work also arises as a response to the public contest called *Traveling Salesman Challenge* [**?**], issued by *Kiwi*, a well established OTA, even though the beginning of this work dates prior to the issue of the challenge. In this challenge, Kiwi recognizes that, in most cases, users do not care about the order in which they visit a given list of cities, and that there exists a market niche interested in this type of services. Kiwi also recognizes that most OTA do not offer these type of services due to the computational complexity associated with the problem.

This work intends to address the problem of solving multi-city trips, by studying it, and by developing the necessary technologies to effectively solve the problem in a time-efficient manner. It also wishes to develop a proof of concept online flight search application, implementing these technologies in order to, ultimately, provide high-quality search for multi-city trips.
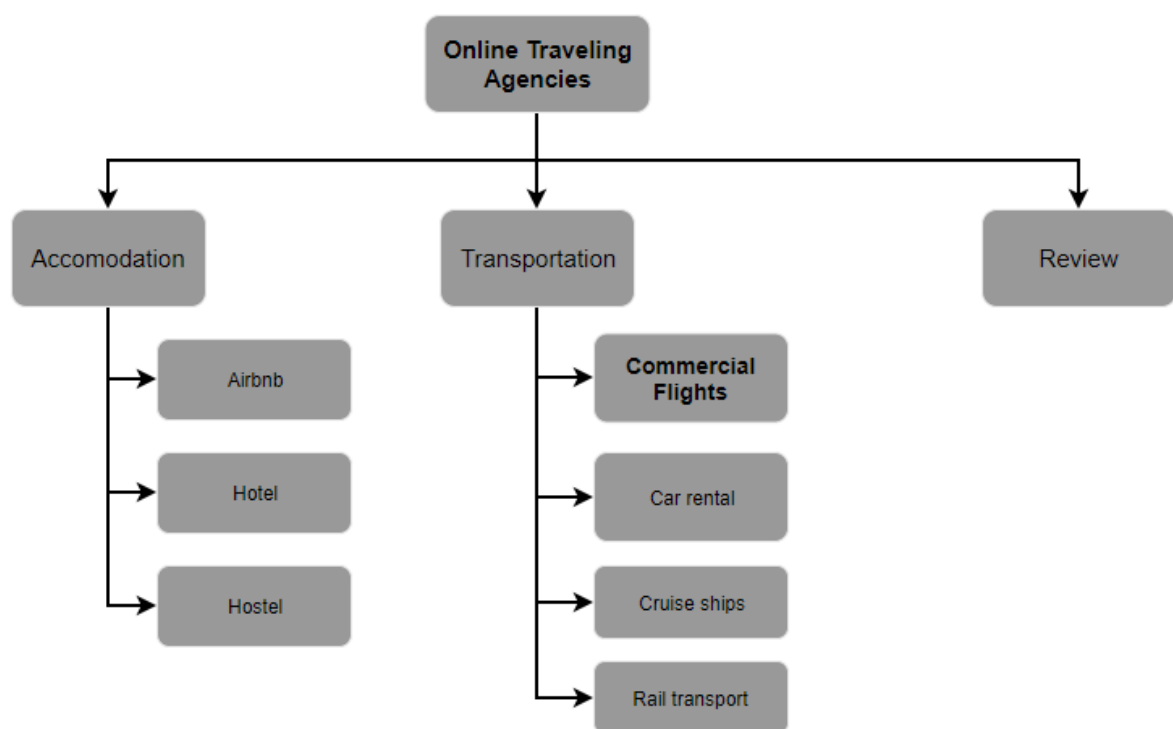
1.2   Existing Services

The tourism industry exists since, at least, the 19th century, but it was impacted by some significant chapters of human technology, which lead to an increased and sustained growth of the market size. First, during the 1920's, the development of commercial aviation meant a significant impact on the industry, shifting the transportation focus to the airplane. Much later, during the 90's, the establishment of the internet led to some changes in the market because airlines were able to sell directly to the passengers [**?**]. More recently, the widespread use of mobile phones and lead to new increases in the market size. In 2016, the direct contribution of the tourism industry for the GDP was over 2.3 trillion dollars, while the total contribution was over 7.6 trillion [**?**].

The increase in the market size of the tourism sector is sustained by traveling agencies, whose main function is to serve as an agent, advertising and selling products and services on behalf of other service providers. These services usually include, but are not limited to, transportation, accommodation, insurance, tours and other tourism associated products. There are both online and offline travel agencies, and the focus of this chapter will be given to the online traveling agencies.

Online Traveling Agencies are websites and applications which offer traveling services or reviews, as illustrated in figure **??**. Most OTA function as a metasearch engine, which means that they search multiple independent travel services providers. Furthermore, this is the main differentiator from the search engines provided by OTA's, and those from direct travel suppliers, as are the websites of individual airlines. Direct travel suppliers are limited to display the results of their own services. On the other hand, OTA usually do not own any travel services but serve solely as an intermediary between traveler and travel services provider. In some cases, these metasearch engines may resort to web scraping in order to get real-time data on the flights provided by different airlines. Recent reports show that while OTA are increasing its market share, direct travel suppliers still account for 57% of the total online travel consumption [**?**].

In order to better understand the difference between OTA's and direct travel suppliers, consider the case of a user searching for a simple round flight between two cities. If this user visits an individual airline website, the flight results presented are limited to those offered by this airline. However, there is no guarantee that the airline flies the user defined route. Furthermore, even if there is such a route, it is possible that this route is not flown every day. Note that these two types of problems are very common, especially in low-cost airlines. In contrast, the same round flight search on a metasearch engine of some OTA will produce a variety of results which include several different airlines. Finally, since OTA aggregate data from different airlines and other meta searches, it is less likely for the two problems described above to occur using the OTA search tools. In conclusion, collecting data from multiple sources usually results in a higher variety and quality of results.

Despite the variety of services provided by OTA, this section will focus only on making a review of

Figure 1.1: The different services provided by Online Traveling Agencies

the existing search tools available for the commercial flight transportation segment. This analysis will be made from both the user and the developer point of view because the offered services often vary according to this.

1.2.1 User Search Tools

Depending on the specific application, online traveling agencies may offer services which include transportation and accommodation. The focus of this section is to create an extensive overview of the search tools available for finding commercial flights. In order to classify the search tools, the results will be grouped according to the *trip type*:

- single/round flight;

- multi-city trip;

This overview will also focus on the different search utilities provided by the search applications. Thus, it will analyze the ability to:

- search over a range of start dates - **DR**;

- present a price overview as a function of time - **PO**;

- present different results for different criteria - **VR**;

- display an interactive map - **Map**;

There are other search utilities, presented by some online search applications, which will not be subject of study. These include:

- price monitoring and alerting service;

- analysis of the probability of price fluctuations;

- information about alternative flights;

- automatic caching for future usage;

This analysis will be considered for the major flight search applications, which include, for example, Google Flights, Momondo, Skyscanner, Kayak, and Kiwi.

Single/Round flights

In order to evaluate the search utilities provided by every flight search application, as well as the actual price displayed, the same search will be performed over all applications. In particular, this search was executed over a month before the actual start date, and corresponds to:

- **S** - single flight - from Lisbon, to Amsterdam, at 8/6/2018;

- **R** - round flight - from Lisbon, to Amsterdam, at 8/6/2018, for a duration of 3 days;

Table **??** displays the results of the search performed over the different applications, together with a checklist specifying which search utilities are provided. It is also worth noting that the price displayed in this table might not correspond to the actual final price. For example, Edreams is well known for charging extra fees only upon the ticket purchase.

Table 1.1: Comparison of the search results, and search utilities, of different online flight search applications, for single and round flights.

company	DR	PO	VR	Map	S	R
Google Flights	✓	✓		✓	65	127
Expedia					81	167
Booking.com / Kayak			✓	✓	67	126
Momondo		✓	✓		62	102
TripAdvisor					56	124
cheapflights.com					75	150
Adioso	✓		✓		69	137
kiwi	✓	✓	✓	✓	76	224
SkyScanner	✓	✓			65	127
Edreams	✓		✓		56	116

From the analysis of table **??**, it is possible to conclude that, for single and round flights:

- different applications provide significantly different results. Single and round flight prices vary up to 44% and 120% respectively;

- half of the applications do not provide date range support;

- half of the applications do not provide price overview as a function of time;

- most applications do not present an interactive map;

Multi-city trips

Following the same methodology as in the previous section, the same multi-city trip search will be performed on the different applications, and the result is presented as **M** in table **??**. This search corresponds to:

6

- single flight from Lisbon to Amsterdam, at 8/6/2018, for a duration of 3 days;

- single flight from Amsterdam to Berlin, at 11/6/2018, for a duration of 3 days;

- single flight from Berlin to Lisbon, at 14/6/2018;

It is worth noting that all flight search applications which currently perform multi-city trips, do it in the exact same way: a multi-city trip corresponds to a collection of single flights, which must specify the origin and destination, together with a single start date.

This case corresponds to a very inefficient search for a variety of reasons. First of all, it constrains the search to one particular route. Second, it does not allow any kind of flexibility around the start date or duration of stay associated to each city. And finally, changing one parameter of the search might affect the rest of the trip.

Once again, the results of the search, together with the search utilities, are presented in table **??**. Note that MTCS stands for the ability to search for multi-city trips, as those specified above.

Table 1.2: Comparison of the search results, and search utilities, of different online flight search applications, for multi-city trips.

company	MTCS	DR	PO	VR	Map	M
Google Flights	✓					184
Expedia	✓					420
Booking.com / Kayak	✓					191
Momondo	✓			✓		178
TripAdvisor	✓			✓		238
cheapflights.com						
Adioso						
Kiwi	✓		✓			226
SkyScanner	✓			✓		238
Edreams	✓					333

From the analysis of table **??**, it is possible to conclude that, for multi-city trips:

- different applications provide significantly different results, which vary up to 135%;

- most applications provide (limited) multi-city trip search;

- most applications do not present different results for different search criteria;

- no application offers the ability to perform multi-city trips on a range of start dates;

- no applications displays an interactive map;

1.2.2 Developer Search Tools

For a developer that intends to create a flight meta-search engine, a way to access flight data is essential. There are some online traveling agencies which extend their API for public consultation. Amongst the

services that these public API's offer, are the possibility of searching for cached, and sometimes, real-time flight data. In some cases, these API's extend their range of services and include endpoints for the consultation of hotel information, car-rental, cruise ships and even railroad data [**?**, **?**].

Usually, these type of content provider benefits from sharing their API, because it generally increases their traffic and the number of potential clients and bookings. In most cases, metasearch engines do not handle booking directly, but redirect to the original content provider. In this case, the content provider retains the entirety of the booking fees. In the case where the meta-search engine does manage the booking, the content provider retains the majority of the booking fee's, while the second party may retain some share of these revenues [**?**].

Amongst the companies which extend their API for third party are Google, Skyscanner, Expedia and Kiwi. These API usually operate as a *free*, *limited* or *private* service. For example, while Expedia and Kiwi offer free publicly accessible API's, Skyscanner has a *private companies only* [**?**] policy, denying any free public consultation for educational purposes. In its turn, Google has a free, but limited API, in the sense that only 50 free daily queries are available, and anything beyond that has an associated cost [**?**].

Due to the educational purpose of this work, the interest in a flight API is limited to those which are free. Considering the different API's that were consulated and analyzed, Kiwi's free public API stroke as the most useful, efficient and fast. The major search utilities that this API offers are [**?**]:

- access to relevant information about cities, airport data and airline companies;

- access to flight information, with flexible queries;

- the possibility of specifying multiple start dates;

- the possibility of specifying a variable duration associated to a return flight;

- flexibility on the specification of the origin and destination;

- possibility to produce multi-flight queries, by grouping up to a total of 9 simple flight queries.

Note that although Kiwi, and other companies, enable the possibility of quering for multiple flights at once, each flight must specify a particular pair of cities and date. Thus, this is a constrained multi-city search, and does not actually give any information about the best route, schedule or set of flights for a multi-city trip.

1.3 Objectives

This work intends to study the time-dependent traveling salesman problem and develop an algorithm to present high quality solutions in a short period of time. This algorithm is to be used in the development of

an online flight search application capable of finding the best flights for a trip composed of several cities, with no particular routing constraints. The concretization of this search engine requires a user interface to collect trip requests, and both access to real time flight data, and to an optimization algorithm capable of solving the TDTSP instance.

Furthermore, a secondary objective is to develop a search engine capable of solving multi-objective tsp instances. This is believed to be important, because many users care not only about the flight price, but other attributes, as the flight duration.

Finally, the search options of the developed search engine should include the possibility to constrain one particular city to a specific time window. The objective of this is to study the possibility of using the developed algorithm to solve constrained problems, as is the TDTSP with time windows.

1.4 Document structure

Here we describe the document structure

2

Literature review

Contents

This chapter is a review on the literature around Combinatorial Optimization, Traveling Salesman Problem and its time-dependent variation, and the state of the art of the algorithms which address these problems.

2.1 Combinatorial Optimization and Routing Problems

2.1.1 Combinatorial Optimization and Computational Complexity

Lawlers, (1976), [**?**], defined combinatorial analysis as "the mathematical study of the arrangement, grouping, ordering, or selection of discrete objects, usually finite in number". Schrijver (2002), [**?**], following this definition of Lawler, improves it with the important concept of optimal solution, "Combinatorial optimization searches for an optimum object in a finite collection of objects.". This definition is followed by a remark, stating that "typically, the collection (...) grows exponentially in the size of the representation", and concluding that "scanning all objects one by one and selecting the best is not an option".

Following a more consice definition [**?**], a combinatorial optimization problem can be defined as follows:

Definition 1) A combinatorial optimization model $P = (S, \Omega, f)$ consists of:

1. a search space S, defined by a finite set of decision variables, each with a domain;

2. a set Ω of constraints amongst the decision variables;

3. an objective function $f \colon S \to \mathbb{R}_0^+$, to be minimized.

The search space is defined by a set of decision variables $X_i, i = (1, ..., n)$, each associated to domain D_i, which specifies the possible value of each decision variable. An instantiation of a variable is an assigment of a value $v_i^j \in D_i$ to a variable X_i. This leads to the definion of a feasible solution $s \in S$, which corresponds to the assignment of a value to each decision variable, according to its domain, in such a way that all constraints in Ω are satisfied. Finally, the objective of the problem is to find a global minimum of P is a solution $s^* \in S$, such that $f(s^*) < f(s) \forall s \in S$, and the set of all set of all global minima is denoted by $S^* \subseteq S$.

When working on a combinatorial optimization problems, it is usefull to have an idea of how difficult the problem is. This characterization is provided by a field called computation complexity. A combinatorial optimization problem Π is said to have worst-case time complexicity $O(g(n))$, if the best algorithm for solving Π finds an optimal solution to any instance of size n of Π, in a computation time upper bounded by $g(n)$.

A problem Π is said to be solvable in polynomial time if the maximum amount of computing time necessary to solve any instance of size n is bounded by a polynomial in n. If k is the largest exponent of such a polynomial, then the combinatorial optimization problem is said to be solvable in $\mathcal{O}(n^k)$ time.

A *polynomial time algorithm* is characterized by a computation time bounded by $\mathcal{O}(p(n))$, for some polynomial function p, where n is the size of the problem instance. If k is the largest exponent of such a polynomial, the problem is said to be solvable in $\mathcal{O}(n^k)$. On the contraty, any algorithm whose computational time can not be bounded by a polynomial function is refered to as an *exponential time algorithm*. Any problem that can be solved in polynomial time is said to be *tractable*, while problems that are not solvable in polynomial time are called *untractable*.

In the field of computational complexity, there is an important concept called *polynomial time reductions*, which transform a problem into another problem, in polynomial time. If the latter problem is solvable in polynomial time, so is the first one. The class of problems which is solvable in polynomial time is called P. On the other hand, there is a class of problems called NP, which stands for *non-deterministic polynomial acceptable problems*, for which given a solution can be *verified* in polynomial time. The class of NP-complete refers to the most difficult problems in NP.

It is worth mentioning that there exists another class, called NP-hard, for which each problem is as hard as the hardest NP-complete problem. More precisely, a problem H is NP-hard when every problem in NP can be reduced to H using a polynomial time transformation. This definition of the class NP-hard leads to the logical conclusion that finding a polynomial time algorithm for *any* problem in NP-hard, would imply the resolution of *all* NP-complete problems in polynomial time. However, up untill now, no polynomial time algorithm was found for any NP-hard problem. Note that the class NP-hard does not necessary belong to the class NP, but this is used following the name convention.

There exists an infamous discussion amongst the cientific community regarding the question "$P = NP?$", since it is one of the major unsolved problems in computer science. We present image (**??**) illustrating the classes according to both possible solutions to the aforementioned question.

2.1.2 The Traveling Salesman Problem

Given a list of cities and the distances between them, the Traveling Salesman is the combinatorial optimization problem of finding a minimum length route which connects every city. By this original definition, proposed by [**?**], the focus of the TSP is to perform optimization on routing problems, as the school bus problem, studied by Merrill Flood in 1942, [**?**], minimizing the total distance of a tour. Some variations of the original formulation allow the adaptation of the problem to suit different optimization goals, [**?**]. For example, instead of distance, the focus may be the minimization of the total cost, travel time, or some other attribute associated to the problem under consideration. It is also possible to search for a route which minimizes two, or more, objective functions at once, [**?**]. In some routing problems, the tour under consideration must satisfy some constraints, [**?**]. Most often, these constraints refer to scheduling conflicts which must be satisfied, [**?**]. A practical application of this is the resolution of routing problems with time windows, [**?**].

NP-Hard

NP-Complete

NP

P

Complexity

$P \neq NP$

NP-Hard

P = NP =
NP-Complete

$P = NP$
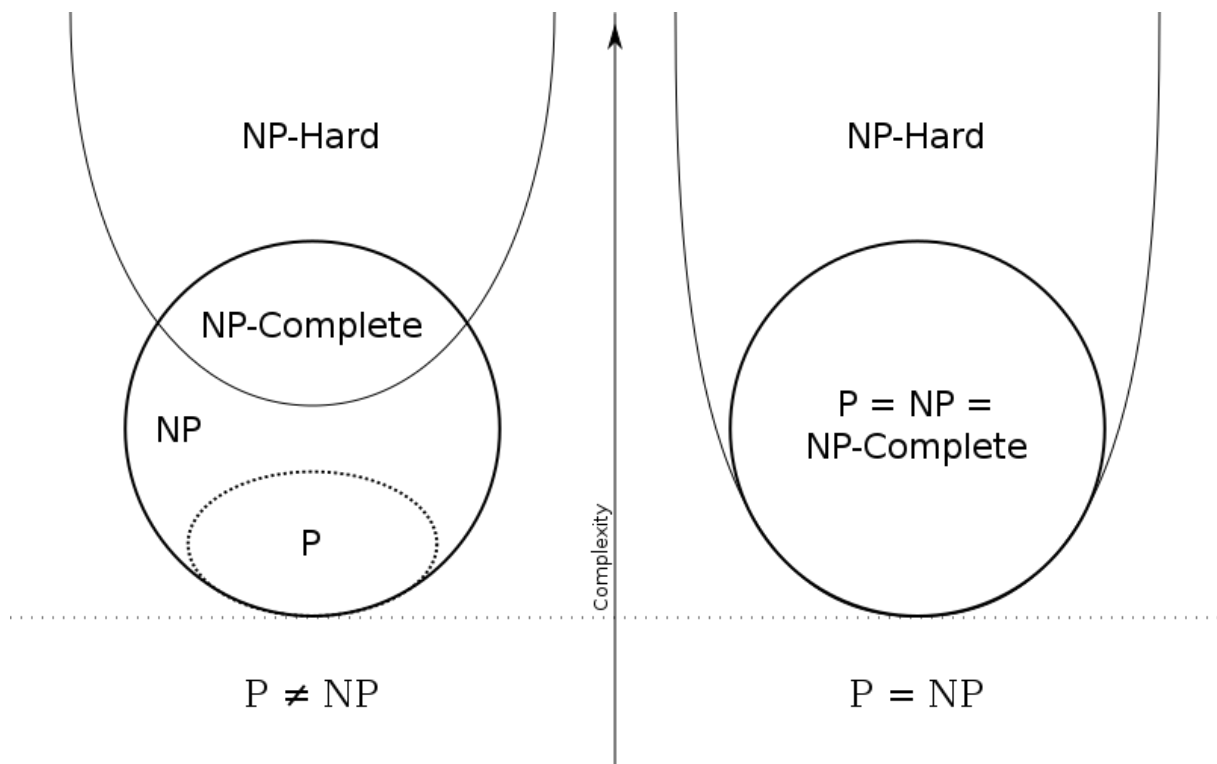
Figure 2.1: Illustration of the classes P, NP, NP-complete and NP-hard

In the field of Graph theory, the TSP is the problem of finding a minimum cost Hamiltonian cycle over a complete, undirected, weighted graph, [?]. The problem of finding a minimum cost Hamiltonian cycle was shown to be NP-complete. This implies the NP-hardness of the TSP. [?]. In several graph problems, considering a symmetric cost between two points is not suitable. This is known as the asymmetric TSP, and it considers a directed graph instead, [?].

In the flight industry, the Traveling Salesman has vast applications. It was applied to airport scheduling, [?]. More recently, the TSP and its time dependent variation have been focus of attention in fields related to Unmanned Aerial Vehicle routing, ([?], [?]). There is also an online website which introduces the Air Traveling Salesman, which introduces the problem of finding layover airports when no direct route is available, [?].

In some cases, the classic formulation of the Traveling Salesman does not adequately describe the characteristics of the problem under consideration. To overcome this, different problem formulations are considered. An example of this is the time dependent TSP, [?]. In this formulation, the cost of each arc is not constant, but varies as a function of time. In general, this problem is harder to solve, [?]. There are several other combinatorial optimization problems which benefit from considering a time dependent approach, [?]. The Vehicle Routing Problem is a field which particularly focus on this problem, due to the characteristics of street traffic, [?]. It is worth mentioning that the Traveling Salesman Problem is a

special case of the Vehicle Routing Problem, in which the fleet is composed by only one vehicle, the salesman, [**?**]. Because of this, works related to the VRP are also interesting for the resolution of the TSP.

Most people are faced with similar routing problems every day. Consider the problem of walking or driving from point A to point B. This is a graph problem, in which the arcs are the streets, the nodes the streets intersections, [**?**]. In its turn, the weights refer to the distance or travel time, which in its turn may be affected by other parameter, as traffic, [**?**]. If the person is familiar with the graph, they are capable of finding a good route mentally, very fast, [**?**]. If the undertaken route is to visit a set of points exactly once, before returning to the original starting point, this is known as the Traveling Salesman Problem.

2.1.2.A Problem definition

The Traveling Salesman occurs both as combinatorial optimization and as a Graph Problem. In either case, the TSP is defined by a graph $G = (N, A)$, where N is the set of nodes, and A is the set of arcs connecting those nodes. The set of nodes is of size $n = |N|$, while the size of the set of arcs is $m = |N|^2$. Each arc, $a_{ij}, i, j \in N, i \neq j$ has an associated weight, c_{ij}, which represents, for example, the distance between cities. The set of arcs is fully connected, that is, each node is capable of reaching any other node directly, without visiting a third node. When two nodes can not be connected by an arc, the cost of that node is considered as very high. In the classical TSP formulation, the graph is undirected, which implies symmetry in the costs of the arcs, that is: $c_{ij} = c_{ji} \forall i, j \in N$. Because of the characteristics of this TSP formulation, the graph is said to be connected, weighted and undirected.

The objective of the TSP is to find the minimum cost Hamiltonian cycle, that is a path which visits each node exactly once, and returns to the initial node, closing the path. A generic solution to the TSP is any permutation σ over the set of nodes, N. σ is also a set, where $\sigma_i, i \in len(\sigma)$, represents the node in the i'th index of the cycle. The cost of a cycle is given by the sum of the weights of each arc by which it is composed, that is $C(\sigma) = \sum_{i=1}^{n} c_{\sigma_i \sigma_{i+1}}$, where $\sigma_{n+1} = \sigma_1$

The TSP can solve diferent types of problems by optimising different parameters. In the classic formulation, the weight between of an arc connecting two nodes represent the distance between two cities. However, the weight of an arc can represent different things, particularly, travel time or travel cost. Althoug changing the parameter may lead the TSP formulation intact, in some cases, it changes the problem. This occurs for example when considering that costs are time-dependent and this will be approached in section **??**.

2.1.2.B Common TSP variations

The majority of the problems which are variations of the classic TSP, have the problem structure altered by some differences concerning the characteristics of the arcs or arcs costs, as occurs with the assy-

metric and metric TSP. In other cases, the variation of the problem may refer to constraints amongst the variables as occurs in the TSP with time windows described in the following subsection, or in the objective of the optimization, as occurs with the bottleneck TSP. The definitions provided in this section are with respect to those provided by [**?**].

Assymetric TSP

In the assymetrical TSP, the cost matrix is not symmetric. That is, there is no constraint imposing that $c_{ij} = c_{ji}$, as happens with the classical TSP.

The ATSP may be more adequate than the TSP for some specific real world problems. For example, when considering a routing problem over a city, some roads may not be connected in both ways. In this case, the weight of an arc connecting two points is different, depending on the direction of traversal of the arc.

Metric TSP

The metric TSP is a special case of the TSP, in which the arcs cost, in addition to being symmetric, also respect the triangle inequality. That is, $c_{ij} \leq c_{ik} + c_{kj}, \forall \, i, j, k \in N$.

Euclidean TSP

In the Euclidean TSP, the set of nodes is placed in a d-dimensional space, and the weight of each arc is given by the euclidean distance. This distances is calculated based on equation **??**, for two points $x = (x_1, x_2, ..., x_d)$ and $y = (y_1, y_2, ..., y_d)$.

$$\left(\sum_{i=1}^{d} (x_i - y_i)^2 \right)^{1/2} \tag{2.1}$$

The euclidean TSP is a variation which is both symmetric and metric.

Bottleneck TSP

In the Bottleneck TSP, the objective is to find a valid route which minimizes the cost of the highest cost arc of the tour. According to the characteristic of the graph, the Bottleneck TSP may either be symetric, assymetric, metric or time-dependent.

The Messenger Problem

The Messenger problem, also known as the wondering traveling salesman, is the problem of finding a minimum cost hamiltonian path connecting edges u and v of the graph G. It can be seen as a Traveling Salesman Problem in which the tour is not closed, but ends on a specific node, different from the initial one. The Messenger problem can be transformed into the TSP, by considering a cost of $-M$ for the arc (v, u), where M is a large number. If the nodes u and v are not specified, and one wishes to find a minimum cost hamiltonian path in G, this can be achieved by a graph transformation, adding one node and connecting it to all other nodes by arcs of cost $-M$. The optimal solution to the TSP on this modified graph can be used to produce the optimal solution to the original problem.

Generalized TSP

In the Generalized TSP, the set of nodes is partitioned into k clusters $V_1, V_2, ..., V_k$, and the objective is to find a shortest cycle which passes through exactly one node from each cluster V_i for $1 \leq i \leq k$. If the dimension of each cluster is 1, that is, if $|V_i| = 1$ for all i, the problem reduces to the TSP. There exists effective graph transformation techniques which reduce the GTSP into the TSP. The GTSP has interesting applications to the tourism industry. For example, a person may want make a world trip and visit one city in each continent. In this case, the problem can be stated as a GTSP instance, in which the clusters are the continents.

The *m*-salesmen problem

In the *m*-salesmen problem, there are m salesmen positioned in node 1 of G. Each salesman visits a subset X_i of nodes of G exactly once, starting and returning to node 1. The objective is to find a partition $X_1, X_2, ..., X_m$ of $V - \{1\}$, and a route for each salesman such that:

1. $|X_i| \geq 1$ for each i;

2. $\cup_{i=1}^m X_i = V - 1$;

3. $X_i \cap X_j = \emptyset$;

4. the total distance travelled by all salesman is minimized.

2.1.2.C   Time Dependent TSP

The time-dependent traveling salesman problem is a generalization of the TSP where arc costs depend on their position in the tour, [**?**, **?**]. This section first introduces the TDTSP as a graph problem, followed by the definition as a sequencing problem.

TDTSP as a graph problem

Let $N = 1, 2, ..., n$ and let $N_0 = N \cup \{0\}$. The TDTSP on a complete graph $K(N_0)$ can be modeled as an optimization problem over a layered graph (V, A). V is the set composed by the source node 0, the termination node T, and intermediate nodes (i, t) for $i, t \in N$. In this representation of the intermediate nodes, the first index of (i, t) identifies the vertex i of the graph $K(N)$, and the second index represents the position of vertex i in a path between nodes 0 and T. In its turn, A is the set of arcs connecting each node. This set is composed of initiation, intermediate, and termination arcs. For $i \in N$, $(0, i, 0)$ denotes an initiation arc from node 0 to node $(i, 1)$, and (i, T, n) denotes a termination arc from node (i, n) to node T. Given $i, j \in N$ such that $i \neq j$, and $1 \leq t \leq n - 1$, (i, j, t) denotes an intermediate arc from node (i, n) to node $(j, t + 1)$. The third index of an arc represents its layer, that is, the position in which it occurs in the path, or in other words, the time of the arc traversal, if we consider 1 time unit for each of these traversals.

When working on the TDTSP, it is often convenient to define $G(n)$ as a subgraph of (V, A), induced by $V \backslash \{0, T\}$. This way, $G(n)$ has n^2 nodes $(i, t) : i, t \in N$ and all the $n(n - 1)^2$ intermediate arcs of A.

A path with n vertices in $G(n)$ is of the form $(v_t.t) : v_t \in N, 1 \le t \le n$. Since consecutives nodes are in consecutive layers, the path can be described by an ordered array $(v_t : t \in N)$. This path can be extended to a $0 - T$ path of (V, A) by appending node 0 and T to the begining and end of the tour, respectively.

The classical TSP and its time-dependent variation share the same objective, which is to find the minimum cost hamiltionan cycle over graph (V, A). Another property they share is the possibility of working over a symetric or assymetric problem.

TDTSP as one-machine sequencing problem

In operation scheduling, the time dependent problem TSP can also be stated as a one-machine sequencing problem, [**?**].

Consider a set of n jobs, $J_1, ..., J_n$, to be executed on a single machine. Each job has a setup cost $C_{i,j}^t$, ocurring when job J_i, processed in the t-th time unit, is followed by job J_j, processed in the (t+1)-th time unit. Consider that each job completion takes exactly one time unit. The machine is in some given initial state, denoted by 0, before the job processing beings. As happens with the classical TSP, the machine has to be returned to its original state, after the job processing ends. The problem is to find a sequence, $J_{w(1)}, ..., J_{w(n)}$, that minimized the total set-up cost C(w), defined by:

$$C(w) = C_{0,w(1)}^0 + \sum_{i=1}^{n-1} C_{w(i),w(i+1)}^i + C_{w(n),0}^n \tag{2.2}$$

It is important to note some important characteristics of this formulation. First, problems with unspecified initial/final state can be formulated in the same way using 0 as the initiation/termination cost, that is $C_{0,w(1)}^0 = 0$ and $C_{0,w(1)}^n = 0$. Secondly, the overwriting of the above formulation reduces the defined problem into the classical TSP and the classical Assignment Problem. The first case is achieved by considering that setup costs are not time-dependent, that is, $C_{i,j}^t = C_{i,j}$. The latter is accomplished by considering that setup costs are not dependent on the second/first job, that is, $C_{i,j}^t = C_i^t$.

2.1.2.D   TSP with time-windows

The Traveling Salesman Problem with Time Windows (TSPTW) is a generalization of the TSP, in which the objective is to find a minimum cost hamiltionian cycle which visits every city in its requested time window. This problem has important applications in the field of routing and scheduling, and is particularly relevant for the Vehicle Routing Problem, as it may occur as a real world constraint imposed by customers, whose operation hours are limited to a time window, [**?**]. Being a generalization of the TSP, the TSPTW is also NP-complete, [**?**]. This section introduce two definitions, one for the asymetric TSP with time-windows, and a second one for the time-dependent variation of this problem.

Below we present a formal definition of the assymetric traveling salesman problem with time windows

(ATSPTW), based on [**?**]. Consider a complete digraph $G = (V, A)$, where V is the set of $n = |N|$ nodes, and A the set of arcs, each associated with a nonnegative arc cost, $c_{i,j}$, and nonnegative *setup times*, t_{ij} associated to each arc $(i, j) \in A$. The nodes correspond to jobs to be processed (as described in the single-machine sequencing problem), and arcs correspond to job transitions, where the setup times, t_{ij}, defines the changeover time needed to proccess node j right after node i. Each node $i \in A$ has an associated *processing time* $p_i \geq 0$, a *release date* $r_i \geq 0$ and a *deadline* $d_i \geq 0$, where the release date and deadline denote, respectively, the earliest and latest possible starting time for the processing of node i. The *minimal time delay* for processing node j immediatly after node i is given by $v_{ij} = p_i + t_{ij}$. The interval $[r_i, d_i]$ is called the *time window* for node i. The time-window is said to be relaxed if $r_i = 0$ and $d_i \to +\infty$. On the contrary, a time-window is called active if $r_i > 0$ and $d_i < +\infty$. It is possible to reach a node $i \in A$ at a time $t \in \mathcal{Z}^+ \cup \{0\}$, sooner than its release date a_i. In this case, it will be undergo a waiting time $a_i - t$, before leaving node i at time a_i.

When dealing with routing problems with time-windows, it is often necessary to define if the time-windows are *hard* or *soft* constraints. Hard time windows consider that a node $i \in A$ can not be visited after its deadline d_i. On the contrary, considering soft constraints, the node i might be visited after the deadline d_i, but in this case a penalty occurs.

The objective function of the problem under consideration depends on specific definition of the problem, particularly, according to the time-window constraints. Dealing with hard constraints, the objective function is defined by the sum of the costs of each arc belonging to that tour, while using soft constraints, the objective function depends on the specific problem definition and the values associated to the aformentioned penalties.

There are several versions of the TSPTW which introduce time-dependent variations. These variations usually focus on time-dependent arc costs, setup times, or processing time. This generalization may occur, for example, as a result of considering the traffic effects associated to real world routing problems. Below is the formal definition of the ATSPTW with time-dependent travel times and costs (ATSPTW-TDC), as defined in [**?**].

Let $G = (V, A)$ be a simple directed graph, $V = \{v_i\}_{i=0}^n$ its set of vertices, where v_0 is the depot vertex. Each vertex v_i has an associated time window $[a_i, b_i]$, verifying that $a_i, b_i \in \mathcal{Z}^+ \cup \{0\}$ and $[a_i, b_i] \subseteq [a_0, b_0] \forall i \in \{1, ..., n\}$. Every time window $[a_i, b_i]$ has associated $p_i = b_i - a_i$ instants of time $\{a_i + k - 1\}_{k=1}^{p_i}$. For simplicity, we will denote $t_i^k = a_i + k - 1$, and therefore $t_i^k \in \mathcal{Z}^+ \cup \{0\}$. The time and the cost of traversing an arc $(v_i, v_j) \in A$ depend on the instant of time t_i^k at which the traversing is started. Consider $c_{ij}^t \geq 0$ and $t_{ij}^t \in \mathcal{Z}^+ \cup \{0\}$, respectively, the cost and the time of travesing an arc (v_i, v_j) starting at instant t_i^k.

The proposed goal in this formulation of the ATSPTW-TDC is to find a Hamiltionan cycle in G, starting and ending at v_0, starting and ending inside the time window $[a_0, b_0]$, such that:

- Starting the circuit at time $t_i^k \geq a_0$ involves a waiting time cost $cwt_0(t_0^k - a_0) \geq 0$ with $cwt_0(0) = 0$;

- The circuit leaves each vertex $v_i \in V$ during its associated time window;

- If the circuit arrives at vertex $v_i \in V$ at time $t \in \mathcal{Z}^+$, such that $t \leq a_i$, it is allowed a waiting time $a_i - t$ with cost $cwt_i(a_i - t) \geq 0$, with $cwt_i(0) = 0$. In this case the circuit leaves vertex i at time a_i;

- The sum of the costs of traversing arcs and of the waiting time costs is to be minimized.

The authors of the work introduced in [**?**], propose an exact algorithm for the previosuly defined ATSPTW-TDC, using several graph transformations, which sucessively reduce the problem into an as-symetric TSP, for which several effecient exact alorithms already exist.

2.1.2.E Multi objective TSP

The multi objective Traveling Salesman problem is a generalization of the classic TSP, and is part of much broader class of problems, the multi objective combinatorial optimization problems, [**?**], and, in particular, multi objective vehicle routing problems, [**?**].

The Multi objective TSP is defined as follows [**?**]:

Given a list of n cities and a set D $= (D_1, D_2, ..., D_k)$ of $n * n$ weight matrix, the objective is to minimize $f(\pi) = (f_1(\pi), f_2(\pi), ..., f_k(\pi))$, with $f_i(\pi) = (\sum_{j=1}^{n-1} d_{\pi(j),\pi(j+1)}^i) + d_{\pi(n),\pi(1)}^i$, where π is a permutation over the set (1, 2, ..., n).

Note that when $D = (D_1)$, this corresponds to single-objective TSP. Note also that the above formulation considers that all objective functions calculate the weight of the hamiltionian cycle, according to the respective weight matrix.

The quality of the results of the multi objective TSP are usually measured according to its performance accross the Pareto criteria, defined as follows [**?**]:

Pareto dominance

A vector $\vec{u} = (u_1, ..., u_k)$ is said to dominate $\vec{v} = (v_1, ..., v_k)$, deonted by $\vec{u} \preceq \vec{v}$, if and only if $\vec{u}$ is partially less than $\vec{v}$, i.e. $\forall i \in \{1, ..., k\} u_i \leq v_i \wedge \exists i \in 1, ..., k : u_i < v_i$.

Pareto Optimality

Pareto optimality is defined asa concept of allocation optimality. An allocation is not pareto optimal if there is at least one alternative allocation which produces improvements

A solution $x \in \Omega$ is said to be Pareto optimal with respect to Ω iff. there is no $x^{'} \in \Omega$ for which $\vec{v} = F(x^{'}) = (f_1(x^{'}), ..., f_k(x^{'}))$ dominates $\vec{u} = F(x) = (f_1(x), ..., f_k(x))$.

Pareto Optimal Set

For a given MOP $F(x)$, the Paretto optimal set $(P*)$ is defined as : $P* = \{x \in \Omega |\ \nexists x^{'} \in \Omega : F(x^{'}) \preceq F(x)\}$

Althoug the above mentioned problem refers to the multi-objective TSP, without loss of generality, the multi-objective optimization can be performed on a time-dependent TSP. There is very few direct research about multi objective Time dependent TSP, but one can cite [**?**], which proposes a multi-objective tabu search for single machine scheduling problems with sequence-dependent setup times.

2.1.3 Vehicle Routing Problem

The Vehicle Routing Problem is the problem of finding the optimal set of routes for a fleet of vehicles, to serve a given set of customers. The VRP is believed to be introduced by Dantzig, in 1959, in a work with the name of *The Truck Dispatching Problem*, in which it is considered a generalization of the Traveling Salesman, [**?**]. It was latter shown that, being a generalization of the TSP, its computational complexity is also NP-hard, [**?**].

Being an NP-hard problem, the focus of the research usually revolves around heuristic algorithms, although there are some procedures which are known to produce optimal solutions, [**?**], [**?**]. As refered by Donati in [**?**], citing the work of Blum, [**?**], even when an exact procedure is available, it usually requires large computational time, which is not viable in the time-scale of hours, as required by this industry.

Malandraki, [**?**], as early as 1992, stated that the assumption of constant and deterministically known costs, is an approximation of the actual conditions of routing problems, and thus, a time-dependent formulation of the problem should be considered. In 1999, Gambardella and colleagues proposed a multi ant colony system for solving the vehicle routing problems using a meta-heuristic approach, [**?**]. Years later, Gambardella expanded this research to include time-dependent variations, [**?**], as proposed by [**?**]. There are several other works, which propose meta-heuristic solutions to solve the time-dependent VRP, [**?**], including the use of simmulated annealing, [**?**], and genetic algorithms, [**?**].

The rest of this section is structured as follows. The next section presents a formal definition of the Vehicle Routing Problem and its time-dependent variation, as well as the most common objectives of the resolution of this problem. Since the TSP occurs only as a generalization of the non-capacitated vehicle routing problem, the study of the capacited vehicle routing is out of the scope of this work.

2.1.3.A Problem definition

Following the definition proposed by [**?**], let $G = (V, A)$ be a graph, where $V = 1, ..., n$ is a set of vertices, representing nodes/customers/cities, with the depot located at vertex 1, and A is the set of arcs fully connecting the nodes. Each arc (i, j), $i \neq j$, is associated with a non negative weight, c_{ij}. Depending on the context of the work, this weight might represent the distance between nodes, the travel time, or even the travel cost. It is assumed that a fleet of m vehicles is available. The Vehicle Routing Problems consists in finding the set of optimal routes such that:

1. each city in $V \backslash \{1\}$ is visited exactly once, by exacly one vehicle;

2. all routes start and finish at the depot;

3. some constraints must be satisfied;

The most common constraints associated to the 4) include: capacity restrictions associated with each vehicle; limit on the number of nodes that each route might visit; total time restrictions; time-windows in which each node must be visited; precendence relations between nodes.

The goal of the Vehicle Routing Problem usually consists in finding an optimal set of routes, as to minimize the total cost, where the cost depends on the total distance covered, and the fixed costs associated to each vehicle. However, depending on the problem under study, the goal may be different, as to minimize the total travel time, minimize the total number of vehicles, or even both at the same time [**?**].

2.1.3.B   Time-dependent VRP

The Vehicle Routing Problem is a very wide class of optimization problems, whose precise problem definition usually depends on the characteristics of the problem under considerations. Thus, introducing time-dependencies on the problem also depend on the specifics of the situation. There are several athors which consider time-dependent travel costs, [**?**], and the objective is to minimize the total costs, while others introduce time-dependent travel times, [**?**], and the objective is to minimize the total travel time. There are also those who consider that the objective function is a function of both travel time and travel costs, and at least one of these (travel time, travel cost) is time-dependent, [**?**]. The definition here proposed follows this last time-dependent variation.

The time-dependent VRP is defined, following the work of [**?**], as follows. Let $G = (V, A)$ be a graph where $A = \{(v_i, v_j) : i \neq j \wedge i, j \in V\}$ is the set of arcs, and $V = (v_0, ..., v_{n+1})$ is the set of vertex Vertices v_0 and v_{n+1} denote the depot at which the vehicles are based. It is considered that each vehicle has an uniform capacity of q_{max}. It is also expected that each vertex in $i \in V$ has an associated demand $q_i \geq 0$, a service time $g_i \geq 0$, and the depot has $q_0 = 0$ and $g_0 = 0$. The set of vertex $C = (v_1, ..., v_n)$ specified the set of n customers. The arrival time of a vehicle at customer i, $i \in C$, is denoted by a_i, and its departure time b_i. Each arc (v_i, v_j) has an associated distance $d_{ij} \geq 0$, and a travel time $t_{ij}(b_i) \geq 0$. Note that the travel time is a function of the departure time from costumer i. The set of available vehicles is denoted by K. Consider that the cost per unit of route *duration* is denoted by c_t, and the cost per unit of route *distance* is denoted by c_d.

In this formulation, there are two goals for the time-dependent VRP. The first corresponds to the minimization of the total number of vehicles used. The second corresponds to the minimization of the total cost, which is a function of both distance and travel time.

There complete definition of the problem follows a mixed integer programming approach, with a total of 11 constraints. These will not be covered in detail here, as the VRP is not the primary object of study of this work, that being the TSP. Thus, it is important to define in which circustances the TDVRP can be transformed into the TDTSP. This is possible by considering only one vehicle, with infinite capacity, and by adapting the objective function according to the problem under consideration.

We conclude this section with the final remark that the above presented definition of the time-dependent VRP corresponds to a static version of the time-dependent case. There is a lot of research around the dynamic case, in which the problem is updated during the execution of the program. This has major applications in the routing industry, and it is often refered to as *real-time* Vehicle Routing. For more information regarding this problem, we refer to [**?**] [**?**].

2.1.3.C Multi-objective VRP

Multi objective optimization corresponds to the resolution of a combinatorial optimization problem in which more than one goals are defined. In the case of the Vehicle Routing Problem, [**?**], the most common objectives include minimizing the fleet size, the total distance traveled, the total time required, the total tour cost, and/or maxizimizing the quality of the service or the profit collected. Note that in most problems, when multiple objectives are identified, the different objectives often conflict with each other.

Multi-objective optimization usually relies on the use of meta-heuristics, [**?**]. There are several works focusing on this problem, and the most promising meta-heuristics for multi-objective optimization include Evolutionary Optimization, [**?**] and Simmulated Annealing, [**?**]. There is also some work considering the Ant Colony Optimization. In particular, a modified ant colony was designed to solve a bi-objective time dependent vehicle routing problem, in which the main goal was the minimization of the fleet size, followed by the minimization of the total cost, [**?**].

2.2 Algorithmic overview

The algorithms which address the Traveling Salesman, and any Combinatorial Optimization problem for that matter, can be classified as exact, heuristic or meta-heuristic. Exact algorithms are thus which always provide an optimal solution to the problem. Although these might seem the first choice, exact algorithms are usually inefficient for solving large problems. In its turn, heuristic algorithms intend to be efficient, abdicating the objective of finding the best solution, and focusing on finding *near-optimal* solutions in a short time. Heuristic algorithms are usually problem specific, while Meta-heuristics algorithms are designed in such a way that they can be used for a variety of Combinatorial Optimization problems, in a fast and efficient way.

2.2.1 Exact algorithms

There are some exact algorithms available for the Traveling Salesman Problem, [**?**], including its time-dependent variations, [**?**], and even with time windows, [**?**]. These algorithms usually require the problem to be formulated as an Integer Linear Programming Instance. In this section we present ILP definitions for both the classical time-dependent TSP. We also present a brienf introduction regarding the Branch and Bound algorithm, which has proven to be very useful for determining exact, or at most, near optimal solutions for the TSP. The software *Concorde*, uses a Branch and Bound algorithm, and was used to solve all 110 instances of the TSPLib, reporting exact solutions in every problem, including a instance with 89.900 nodes, although it required more than 110 CPU years.

2.2.1.A   Integer Linear Programming

ILP for the TDTSP

The TSP be formulated as an integer programming problem, (see Laburthe 1998). The decision variables are x_{ij}, which take values of one and zero, followig the following rule.

$$x_{ij} = \begin{cases} 1, \text{ if the tour contains arc (i,j)} \\ 0, \text{ otherwise} \end{cases} \tag{2.3}$$

Let c_{ij} represent the weight of the arc (i,j). The objective of the TSP is described by equation **??**. Equations **??** and **??** represent constrains over the variables, particularly, that a tour must enter and leave each node exactly once. However, this does not completly define the characteristics of a Hamiltonian cycle. To eliminate the possibility of subtours, that is, of having some node more than once in a solution, it is necessary to introduce the sub-tour elimination constraint. This is expressed in equation **??**. Without this constrain, the formulation of the problem reduces to the classical Asignment Problem, that can be solved in polynomial time, $\mathcal{O}(n^3)$.

$$min \sum_{ij} c_{ij} x_{ij} \tag{2.4}$$

$$\forall i, \sum_{j} x_{ij} = 1 \tag{2.5}$$

$$\forall j, \sum_{i} x_{ij} = 1 \tag{2.6}$$

$$\forall S \subset N, S \neq \emptyset, \sum_{i \in S} \sum_{j \notin S} xij \geq 2 \tag{2.7}$$

ILP for the TDTSP

The TDTSP can be formulated as integer linear programming problem, by using bynary decision variables, x_{ijt}. These variables take a value of zero or one, according to the rule of equation **??**.

$$x_{ijt} = \begin{cases} 1, \text{if city j and i are visited in the time period t and t-1, repectively} \\ 0, \text{otherwise} \end{cases} \tag{2.8}$$

The objective function is presented in equation **??**. Equations **??**, **??** and **??** represent constraints over the decision variable. Particularly, they state that each city must be entereded exactly once, left exactly once, and visited in exactly one time period, respectively. As occurs with the classical TSP, the ILP formulation needs to formulate a constraint to eliminate the possible formation of sub tours. This is presented in equation **??**. Finaly, equation **??** guarantees that the decision variable takes binary values.

$$min \quad \sum_i \sum_j \sum_t C_{ijt} x_{ijt} \tag{2.9}$$

$$\sum_j \sum_t x_{ijt} = 1 \quad i = 1, ..., n \tag{2.10}$$

$$\sum_i \sum_t x_{ijt} = 1 \quad j = 1, ..., n \tag{2.11}$$

$$\sum_i \sum_j x_{ijt} = 1 \quad t = 1, ..., n \tag{2.12}$$

$$\sum_{j=1}^{n} \sum_{t=2}^{n} t x_{ijt} - \sum_{j=1}^{n} \sum_{t=1}^{n} t x_{ijt} = 1 \quad i = 1, ..., n \tag{2.13}$$

$$x_{ijt} \in 0, 1 \quad i, j, t \in i, ..., n \tag{2.14}$$

2.2.1.B Branch and Bound

Branch and Bound ($B\&B$) is one of the most used tools to solve large NP-hard combinatorial optimization problems. To be precise, $B\&B$ should be classified as an algorithm paradigm, constitued by 3 main parts, which have to be choosen according to the problem under consideration, and for which many options may exist, [**?**].

The force of the $B\&B$ comes from it being a search algorithm which (indirectly) searches the complete search space of the problem. Since this is not directly feasible, due to the common exponential growth of the solution space, $B\&B$ takes advantage of *bounds*, combined the information about the current best solution, to safely discard certain solutions amongst the search space.

At any point of the algorithm, there is a *current solution*, and a *pool* of unexplored subsets of the

25

solution space. At the beginning of the algorithm, this pool consists of (only) the root node, and at the end of the algorithm, it will consists of an empty set, meaning that the entire search space was sucessfully explored. The initialisation of the $B\&B$ requires the *incumbent*, which denotes the objective function value of the current solution, to be initialised as ∞. If it possible to generate an initial feasible solution using some heuristic method, this solution is recorded and its objective value is set as incumbent. The process of generating an initial solution has usually a positive impact on the $B\&B$ algorithm. After the initialisation, this algorithm enters in an iterative proccess, untill the pool of unexplored subsets is empty. This processes consists of three main components:

- selection of a node to process;

- bound calculation;

- branching.

Branch and Bound algorithms vary according to the strategies established for each of the three main components of the iterative proccess, aswell as the initial heuristic. In any case, the bounding function selected is the key for any good branch and bounding algorithm, because the selection of a bad function can not be compensated with good choices on the branching and bounding strategies. For example, consider the trivial case where the bounding function is the constant value of 0. It is obvious that this will always be a lower bound to the problem, but it does not produce any quality information of which solutions to discard. Ideally, the value of the bounding function for a given subproblem should be equal to the value of the beast feasible solution to that problem. This is usually not possible, since subproblems may also be NP-hard. Thus, bounding functions are choosen according to the proximity to the best possible value, and to its time complexity - usually restricted to polynomial time.

To complete this overview, selection strategies for the TSP usually revolve around Best First, Depth First and Breadth First Search. There are several works which discuss the main differences of each selection strategy and report on which strategy might be more adequate according to the problem characteristics. Finally, the branching strategy in the TSP usually consinsts of selecting any node with a degree 3 or higher in the search tree. There are several comments around this strategy, some authors opting for the selection of a node with a low bound, as this theoritically will reduce the number of searches when processing the node with higher bounds.

2.2.2 Heuristic algorithms

In some particular cases, exact algorithms can not be used in the resolution of the TSP problem under question. This usually occurs when dealing with very large instances of the problem, or when there is an urgency in obtaining solutions in a fast way. In these cases, using approximation algorithms may be

a good choice. These algorithms are not guaranteed to produce an optimal solution, however, with a good heuristic, approximation algorithms produce high quality solutions in a reasonable time. Generally, the heuristic may be classified as one of two classes: construction or improvement heuristics, [**?**].

When an optimal solution is not known, measuring the heuristic performance of a method may be difficult. In this cases, the heuristic evaluation can be done by comparison with the Held-Karp lower bound.

Heuristic algorithms are usually specifically designed for a particular optimization problem. For example, the Lin-Kernighan Heuristic was created to solve the symmetric Traveling Salesman Problem, and does not seam to have usefull applications in other combinatorial optimization problems.

2.2.2.A Held-Karp Lower Bound

In some cases, the quality of a heuristic solution can not be directly calculated, as no exact solution for the problem under consideration is known. In this cases, it is important to have a way of evaluating performance. The standart way of doing this is by comparing the heuristic solutions which the solution generated by the Held-Karp (HK) lower bound, [**?**].

The HK lower bound is the solution to the linear programming relaxation of the ILP formulation of the TSP. This solution can be found in polynomial time for moderate instance sizes.. However, for a very large problem, solving the relaxed problem directly is not feasible. In this cases, Held and Karp prose an interative algorithm in order to approximate the solution. This methods involves computing a large number of minimum spanning trees. This iterative version of the algorithm will often keep the solution within 0.01% of the HK lower bound, [**?**].

2.2.2.B Tour construction

A construction algorithm is based on the construction of a valid tour. The construction process stops when a valid tour is found. No improvement over the formulated tour is attempted.

A – Nearest neighbour   The nearest neighbour is a very simple heuristic for the TSP. This algorithm starts with the selection of a random node. Then, while there are unvisited nodes, the heuristic always selects the nearest node which was not yet visited. This proccess is repeated while there are unvisited nodes. Finally, when there are none, the construction is complete with the return to the first node.

The computational complexity of the nearest neighbour is $\mathcal{O}(n^2)$. The solutions generated with this heuristic are often within 25% of the optimal solution.

The pseudocode for the NN algorithm is presented below.

1. Select a random city

2. Select the nearest unvisited node

3. If there are unvisited nodes, repeat step (2)

4. Return to first node

B – Greedy heuristic   The greedy heuristic is a construction algorithm which creates a valid tour by repeatedly selecting the arc with the lowest weights, and taking into account the problems constraints. In particular, the greedy algorithm rejects an arc which creates a cycle with less than n edges, or which would create a sub tour.

The computational complexity of the greedy heuristic is $\mathcal{O}(n^2 log_2(n))$. The solutions generated by this heuristic are often within the 20% of the optimal solution.

1. Sort all arcs according to its weight

2. Select the lowest weight arc, if it does not violate any constraint

3. If the constructed solution is not complete, repeat (2)

2.2.2.C   Tour improvement

Improvement heuristics are algorithms which work over a valid and complete solution in order to improve it. The most commom improvement heuristics are the 2-opt and 3-opt local search algorithms. The Lin-Kernighan algorithm (LKA) is a particular implementation of the above mentioned local searches methods, in which a k-opt local search is employed, where the value of k varies during the algorithm execution. LKA have shown to be very efficient and capable of presenting high quality solutions to the TSP.

A –   2 and 3 opt tours   "In optimization, 2-opt is a simple local search algorithm first proposed by Croes in 1958 for solving the traveling salesman problem."

The 2-opt is possibly the most simple local search algorithm. The objective of this method is to find route crossovers, and fold them. When this occurs, the overal cost of the newly constructed tour will decrease.

The 2-opt search works in a recursive way. This search algorithm tries to improve the original tour, by removing two edges from the original cycle, and reconnecting the two paths created. This process is ilustrated in figure **??**. In a 2-opt search, there is only one way of connecting the nodes in a way which will result in a different and valid tour. If the new tour has a lower cost, the cycle restarts, using the new tour as the improvement object. Otherwise, two other edges are selected, and the cycle continues with

Figure 2.2: The 2-opt local search works by reconnecting two edges, hoping to fold possible crossovers, decreasing the overall tour coust. In the left image, a crossover is identified. In the middle image, a the edges beloning to this crossover are removed, and in the figure to the right, they are reconnected, forming a new valid tour

the original solution. This iterative method continues untill no further improvement is reached over a complete iteration cycle. In this case, the tour is known to be 2-opt.

The 3-opt search is very similar to the 2-opt. Instead of selecting two edges and reconnecting the path, the 3-opt selects 3 edges. In this case, there are two different ways of forming a new valid tour. A 3-opt move can also be seen as two or three 2-opt moves combined in the formation of a new tour. The iterative cycle of the 3-opt search works in the same way as the 2-opt.

B – k opt tour More generally, k-opt local search methods are a way of rearanging a tour, by taking k edges and reconnecting the paths in order to form a new valid tour. Any tour that is known to be $k - opt$ is also $(k-1)opt$. Some particular problems, as "the crossing bridges", figure **??** can only be solved with a 4-opt or higher method.

Figure 2.3: The crossing bridges can only be solved by reordering 4 edges. The resolution of this problem with local search is only possible with 4-opt or higher.

C – Lin-Kernighan heuristic The Lin-Kernighan Heuristic (LKH) [**?**], is an algorithm for the symmetric TSP, which was the state of the art for over than 15 years. LKH is known for producing optimal solutions often, for presenting solutions within 2% of the Held-Karp lower bound, and for having a time complexity of approximately $\mathcal{O}^{n^{2.2}}$, [**?**]. This heuristic is constructed for the symmetric, and using it for the asymmetric generalization requires a graph tranformation process, which transforms the asymmetric instance with n nodes, into an equivalent symmetric one with $2n - 1$ nodes, [**?**] [**?**]. Thus, for the

same number of nodes, solving an assymetric TSP with the LKH heuristic is usually 4 times harder than solving the symmetric case.

To understand the Lin-Kernighan heuristic, it is necessary to think about the TSP in a slighly different manner. Consider the following way of defining a combinatorial optimization problem: "find from a set S a subset T that satisfies some criterion C and minimizes an objective function f." In the TSP, the objective is to find from the set of all edges (S) of a complete graph, the subset (T) which forms a valid tour (C) and minimizes the objective function (f). Using this formulation it is now possible to explain the LK heuristics.

Given an non-optimal and feasible solution T, it is non-optimal because k elements $\{x_1, ..., x_k\}$ in T are *out of place*. To improve this solution, and make it optimal, one would have to substitue the set of k elements $x_1, ..., x_k$ with the elements $y_1, ..., y_k$ of $S \backslash T$. Because there is no knowledge about how many elements are missplaced, Lin and Kernighan consider that setting the value of k a-priori would seem artificial. Thus, they propose an interative procedure in which the algorithm dynamically estimates the best value for k. In order to do this, the LKH first estimates the most out of place elements, x_1 and y_1. Then, with this values set aside, it tries to repeat this proccess for x_2 and y_2, and so on. It stops this inner loop when no improvement seams plausible, replaces the current solution T with the new solution generated from replacing the now selected elements, and restarts the whole proccess. This proccess is formalized below, as presented by Lin and Kernighan.

This heuristic has not been shown to work for the time-dependent TSP, as it is constructed for a symmetrical $n * n$ cost matrix only. Thus, the overview of this algorithm will not be extensive, as it has no pratical application to the problem under consideration. In any case, being a very relevant heuristic for the classical TSP, it is an algorithm worth mentioning.

2.2.3 Meta-Heuristic algorithms

Meta-Heuristic algorithms are heuristic algorithms which, unlike the classical heuristic, can be applied to a variety of optimization problems. Meta-Heuristic are designed to be applied to combinatorial optimization problem, and not to a specific problem of this class. Meta-Heuristic rose in importante during the 1990's, and have become one of the most important class of algorithms in computer science.

More formally, a meta-heuristic is an iterative generation process, which guides a heuristic by combining intelligently different concepts, for exploring and exploiting the search space, using learning strategies to structure information, as to efficiently find optimal or near-optimal solutions, [**?**].

This subsection will introduce a few of the most relevant meta-heuristics in the resolution of the Traveling Salesman Problem, particularly, the Ant Colony Optimization (ACO) and the Simmulated Annealing procedures (SA). There is a variety of meta-heuristics which are not discussed here, but which have also been succesfully applied to the TSP. Examples of these meta-heuristics are the Tabu-Search, Evo-
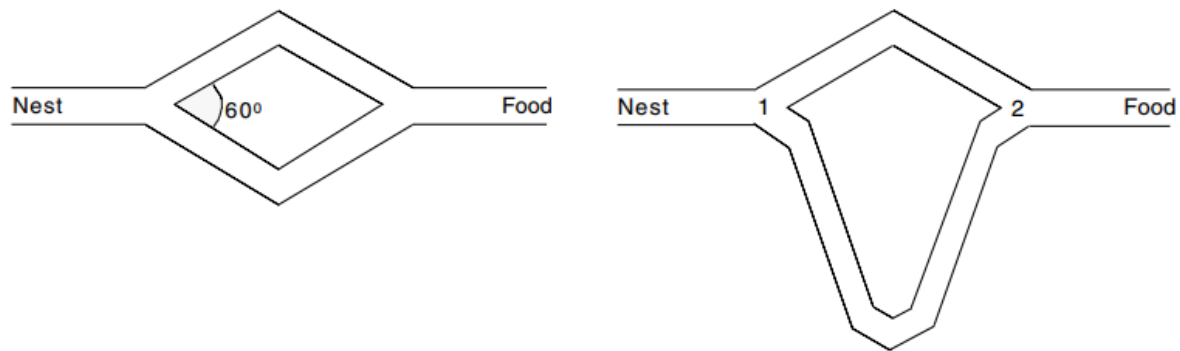
Figure 2.4: The double bridge experiment. On the left, two bridges with the same length. Experimental results show that ants distribute themselves evenly amonst both bridges. On the right, one of the bridges is longer than the other. Experimental results show that ants use the shorter bridge more often.

lutionary Algorithms (EA), in particular the Genetic Algorithm (GA), and many other Swarm Intelligence algorithms, from which the Ant Colony Optimization is the oldest and most widely used.

2.2.3.A Ant Colony Optimization

The Ant Colony Optimization, [**?**, **?**, **?**], is based on the behaviour of real ants, and was developed by M. Dorigo et. al, which were curious about the generelization of the double bridge experiment, [**?**], [**?**], illustrated in (figure **??**). This led to an adaptation of this experiment, substituing the double bridge with a graph, the pheromone trail with artificial pheromone, and the real ants with artificial ants, with presented some extra capabilities intended to facilitate the resolution of more complex problems [**?**].

Using the model of a static combinatorial optimization problem, as defined in section **??**, it is possible to derive a generic pheromone model, that can be exploited by the Ant Colony Optimisation. This means that both the classical and the time-dependent TSP, which can be formulated by the mentioned model, can, potentially, be solved by the ACO algorithm.

The following pseudo-code represents the algorithmic skeleton for the ACO model, and each of its parts will be explained with more detail below.

The general process of the Ant Colony Optimisation algorithms is as follows. The algorithm starts with a parameter initialization. This is also responsible for setting the pheromones levels to some value, τ_0. This value is usually chosen using a heuristic function. For the TSP case, the heuristic chosen is often the nearest neighbour.

Reinsert the below hidden SA metaheristic procedure.

After initialization, and until some specific termination condition is met, the ACO algorithm runs in a loop, which consists of 3 main steps: solution construction, local search (optional), and pheromone update. Each of this steps is detailed below.

The solution construction is a process carried out by each of a specified number of ants. Each ant

starts with an initially empty solution, s_p, and at each iteration step expands its solution with a valid solution component, c_i^j. This construction function differentiates the ACO algorithm for every model. By restricting the construction method to the agents (the ants), the rest of the algorithm does not have to be heavily adapted to the specific model. However, the function that is responsible for selecting the feasible solution components, has to be aware of the model variables and its set of constraints. It has to determine those variables whose addition to the partial solution do not constitute a violation to the set of constraints of the model. This set is represented by $N(s_p)$.

Having the set of all feasible solutions components, $N(s_p)$, it is necessary to choose a single component, c_i^j. This selection is done probabilistically, and the choice takes into account both pheromone (exploitation) and heuristic information (exploration). The algorithmic parameter q_0 is responsible for defining both method's relative importance. The outline execution of the function is presented in equation **??**, and is as follows. A random value, q, is set. If this value is lower than the algorithms parameter q_0, the selection of the node is done with the heuristic rule, equation **??**. Else, the Ant System rule is used, **??**.

$$c_i^j = \begin{cases} \text{heuristic rule,} & if q \leq q_0 \\ \text{ant system rule,} & \text{otherwise} \end{cases} \tag{2.15}$$

$$argmax_{l \in N_i^k} \tau_{il}[\eta_{il}]^\beta \tag{2.16}$$

$$p(c_i^j | s_p) = \frac{\tau_{ij}^\alpha [\eta(c_i^j)]^\beta}{\sum_{c_i^l \in N(s_p)} \tau_{il}^\alpha [\eta(c_{i,l})]^\beta} \tag{2.17}$$

After finalising the construction of valid solutions, the ACO algorithm may implement a local search. Although this step is optional, it has been demonstrated that ACO algorithms reach their best performance when local search is applied. The ant's construction method is biased by the pheromone information, while the pheromones values are biased by the quality of the solutions. Local search, also called Daemon Actions, are techniques which intend to work on the existing solutions, exploring and expanding the search space, and ultimately improving the quality of the solutions. This will be reflected in the ant's construction method, along the next iterations. The most widely used local search methods are the 2-opt search, the 3-opt search, and the Lin-Kernighan heuristic.

The final step of each loop of the algorithm is the pheromones update. This is responsible for making solution components that belong to good solutions more desirable in the following iterations. To achieve this objective, two methods are implemented. The first is pheromone deposition, which increases the pheromone intensity of the solution components belonging to the most promising solutions. The amount of solutions that are used to deposite pheromone is a parameter of the algorithm. The second method to achieve this step's goal, is pheromone evaporation. While it may seam counter-intuitive to deposit

pheromones and, at the same time, also evaporate them, this step is crucial to avoid a rapid convergence to sub-optimal solutions. pheromone deposition alone is responsible for making good solutions more desirable, while pheromone evaporation reduces both the desirability of bad solutions and the sub optimal convergence of good solutions, favoring the exploration of the search space. The pheromone update is implemented as in equation **??**.

$$\tau_{ij} = (1 - \rho)\tau_{ij} + \sum_{s \in S_{upd}|c_i^j \in s} g(s) \qquad (2.18)$$

2.2.3.B Simulated Annealing

The SA algorithm was delevoped using an analogy between the phyisical annealing in solids, and finding the minimum cost configuration in combinatorial optimization problems. In the phyisical world, annealing is the process of heating a metal until the melting point, and reducing the temperature in a controlled way. The decrease in temperatur results in a particle rearangement, in which lower energy states are reached. When the heating temperature is very high, and the temperature is decreased very slowly, this will result in the ground state of the solid - its minimum energy state. The analogy between physical world and the combinatorial optimization problems is achieved by considering that the energy of the metal corresponds to the cost of the solution, and the particle rearangement consists in the selection of a neighbourhood solution, [**?**].

The Simulated Annealing algorithms concists in an iterative improvement, which stochastily accepts up-hill moves. More precisely, the procedure starts with the selection of a feasible solution, aswell as the Initialization of some necessary parameters, as the temperature, which will serve as a control variable. After this, the SA enters an iterative proccess. At the core of this iterative proccess is a local search proccess, which is executed a fixed number of times at each iteration. Authors usually estabilish this value as 2 times the number of nodes. After this local search procedure is complete, the temperature is decreased according to its cooling schedule, after which the local search restarts, and the cycle continues, untill either the execution time is reached, or the temperature reaches 0.

Simulated Annealing differs from other iterative improvement algorithms, because it stochastily accepts up-hill moves, which allow it to escape from local minima (as happens in, f.e., the Tabu Search). More precisely, at each stage of the local search procedure, the difference in the energy level, Δ, between the current state and the newly generated state is calculated. If Δ is negative, the new state is better than the current one, and it is (always) accepted. On the contrary, if Δ is positive, the state is accepted if equation **??** is verified, otherwise, it is rejected. This equation is often called the Metropolis criteria. Note that using this criteria, as the temperature approaches zero, less and less bad states are accepted, and at $T = 0$ Simulated Annealing no detoriations will be accepted at all.

$$exp(-\frac{\Delta}{T}) > Random[0, 1[\qquad\qquad (2.19)$$

Time-dependent scheduling problems can also be solved using this meta-heuristic, [**?**], as the algorithm solely relies on the search of a neighourhood set. The pseudo-code describing the SA procedure is presented below.

In the first works published about the Simulated Annealing, it was proven that if the temperature is cooled very slowly, the process will converge to the optimal solution. More precisely, if temperature drops no more quickly than $C/log(n)$, where C is the Boltzman constant, and n is the number of steps taken so far. This result however is not as relevant as it first seems, because this cooling schedule is *very* slow. Some authors refer that it is faster to do exhaustive search than to follow this colling schedule, [**?**].

The Simulated Annealing procedure varies according to: the cooling schedule; the neighbourhood search criteria; the Markov chain length. There are several reports which describe the influence of these modules in the overall performance of the SA procedure. There are other very relevant aspects of this algorithm which may also influence the results as, for example, the initial and the final temperature.

Reinsert the below hidden SA metaheristic procedure.

3

Problem Formulation

Contents

3.1 Flying Tourist Problem

Consider a tourist who wishes to take a trip that visits every node (city) i in the set of nodes V, $|V| = N$, with no particular order. The start node will be noted as v_0, while the return node as v_{n+1}, and the complete set of nodes is given by $V_c = V \cup \{v_0\} \cup \{v_{n+1}\}$. The trip must start at a time $t \in T_0 = [T_{0m}, T_{0M}]$. Upon visiting a node, the tourist will stay there for a duration of d time-units (days). Consider that for each node to be visited, there is a range for the value d might take, that is $d \in d_i = [d_{im}, d_{iM}]$ and $d_{iM} \geq d_{im} \geq 1$. The complete set of durations associated to each city is given by D, and $|D| = N$. Furthermore, to each city $i \in V$, there is an associated time-window TW, $|TW| = |V| = N$, which defines the set of dates in which the city i may be visited.

By following this definition, the FTP is completely defined by a structure $G = (V_c, A, T_0, D, TW)$, used to create a multipartite graph describing the request. This multipartite graph is divided into k layers, where each layer corresponds to a particular moment in time. Besides this, every node in a layer is connected to all nodes in the subsequent layer. The set of arcs that connects these nodes is given by A. To each arc $a \in A$, it is associated a cost c_a (ticket cost) and a processing time p_a (flight duration), which depend upon the routed nodes, as well as the time in which the arc transition is initiated, that is, $\forall a_{ij}^t \in A$, $c_{ij}^t \geq 0$ and $p_{ij}^t \geq 0$.

A valid solution s to the formulated FTP is a set of arcs (commercial flights) which start from node v_0 during the defined start period, visits every node i in V during its defined time-window $TW(i)$, by considering the staying durations defined by $D(i)$, and finally returns to node v_{n+1}. The set of all valid solutions is given by S. The goal of the FTP is to find the global minimum $s^* \in S$, with respect to the considered objective function.

The objective function associated to this problem depends on the user criteria. While some users might consider the expended cost to be the most important factor, there are others who consider the total flight duration of crucial importance. Thus, a total of three different objective functions shall be herein considered: (i) the expended cost (see eq. **??**), (ii) the flight duration (see eq. **??**), and (iii) the resulting entropy (see eq. **??**), where the latter corresponds to a weighted sum between the former two.

$$F_c(s) = \sum_{n=0}^{N+1} c(s[n]) \tag{3.1}$$

$$F_t(s) = \sum_{n=0}^{N+1} p(s[n]) \tag{3.2}$$

$$F_e(s) = \sum_{n=0}^{N+1} w_c * c(s[n]) + w_p * p(s[n]) \tag{3.3}$$

Figure 3.1: Illustration of a Flying Tourist Problem using a multipartite graph. To each node (A,B,C) it is associated a waiting period of respectively (1,2,3) time units. The red arrows represent a possible solution to the problem.

Figure **??** illustrates the multipartite graph associated to a simple instance of the FTP with $v_{n+1} = v_0 = X$, one possible start date ($t = 0$), 3 nodes to visit (A, B, C), with a fixed duration of respectively (1,2,3) time-units, and no constraints relative to the time-window of each city. A possible solution to this problem instance corresponds to the set of arcs $(a^0_{X,B}, a^2_{B,A}, a^3_{A,C}, a^6_{C,X})$.

Despite the apparent complexity of the proposed definition, it can be used to state very simple flight searches, including one-way and round-trip flights. For example, the problem of finding a single flight from A to B at date T can be instantiated as a FTP given by $v_0 = A$, $v_{n+1} = B$, $T_0 = T$, and $V = D = TW = \{\}$. In its turn, a round-trop flight involving the same two cities and the same start date, in which the staying period in B is b days, is given by $v_0 = v_{n+1} = A$, $T_0 = X$, $V = \{B\}$, $D = \{b\}$ and $TW = \{\}$. Thus, this definition is adequate either for simple and complex trips, which can be customized according to the user search criteria, by setting either an extended start period, or flexible waiting periods.

3.1.1   Relation to the TSP

As previously stated, the proposed FTP is closely related to the TSP and to its time-dependent variation. Given the following list of constraints:

1. $v_{n+1} = v_0$;

2. $T_0 = 0$;

3. $TW(i) = [0, +\infty[, \forall i \in V$;

4. $D(i) = 1, \forall i \in V$;

5. $c^t_{ij} = c_{ij}, \forall i, j \in V, \forall t$.

constraints (1-4) enable the reduction of the devised FTP to a TDTSP, as proposed by J.C.Picard ([**?**]), and the final constraint (5) reduces the problem to the classical TSP.

Since the FTP occurs as a generalization of the TSP, and given that the latter problem is well-known to be Np-hard complex, than so is the former one.

3.1.2 Graph construction

By considering the presented FTP definition, the total number of layers (k) of the devised multipartite graph represents the total time span between the earliest date at which the trip might start and the latest date in which it should finish. The arcs that connect those nodes are divided into three groups: *initial*, *transition* and *final* arcs.

The *initial* arcs are those which might initiate the trip. Consequently, they must start at node v_0, at a time $t \in T_0 = [T_{0m}, T_{0M}]$, connecting v_0 to every node in V. There are a total of $k_i = T_{0M} - T_{0m} + 1$ layers for the initial arcs.

Conversely, the *final* arcs are those that connect every node in V to the return node, v_{n+1}. There are as many final layers as there are initial layers, and the final layer extends from T_{fm} to T_{fM}, where $T_{fm} = T_{0m} + \sum(D)$ and $T_{fM} = T_{0M} + \sum(D)$, where $\sum(D)$ corresponds to the summation of all entries belonging to D. In the example depicted in Figure **??**, there is a single initial and final layer, since there is only one possible start date.

The *transition* arcs are those which fully connect the N nodes belonging to V. The earliest transition arc occurs at a time no sooner than $t_1 = T_{0m} + min(D)$, where $min(D)$ corresponds to the lowest entry of the set of staying durations. Hence, if the trip starts by transiting an initial arc at time T_{0m}, the first transition arc might only be traversed $min(D)$ time-units later. By following a similar approach, the latest transition arc can occur no later than $t_2 = T_{0M} + \sum(D) - min(D)$. Thus, there are a total of $k_2 = t_2 - t_1 + 1$ transition layers, and $k_2 * n * (n - 1)$ transition arcs.

The union of the initial, transition and final arcs gives the set A of all the arcs, which may be used to construct a solution to the requested trip.

Having the information relative to the multipartite graph associated to the devised FTP, it is now possible to construct a three-dimensional array matrix representing this problem, where each entry of the array corresponds to an arc connecting two nodes, at a particular moment in time. This weight matrix is initialized with a very high cost value (as to reject arcs which may not be part of the solution), and every entry of it is updated according to the information of the multipartite graph and the respective objective function. Finally, this weight matrix may be used as input for the optimization system (see section). Insert reference

Although it is clear that any arc $a \in A$ corresponds to a particular flight, it should be noted that no specific or limiting assumption was considered up until now. Instead, it was assumed an entirely abstract arc definition, connecting two nodes at a specific moment in time. In order to transform this set of arcs

into a corresponding set of flights, it is necessary to obtain real-world flight data from some external source. This will be further detailed in section

Insert refer-
ence

3.2   Optimization System

Following the architecture proposed in section **??**, the optimization system will produce a stream of responses to the user request. To produce a fast response, a simple random solution will be generated initially. This is followed with the application of the nearest neighbour heuristic, and concludes with the execution of two meta-heuristic algorithms: the Ant Colony Optimization and Simmulated Annealing.
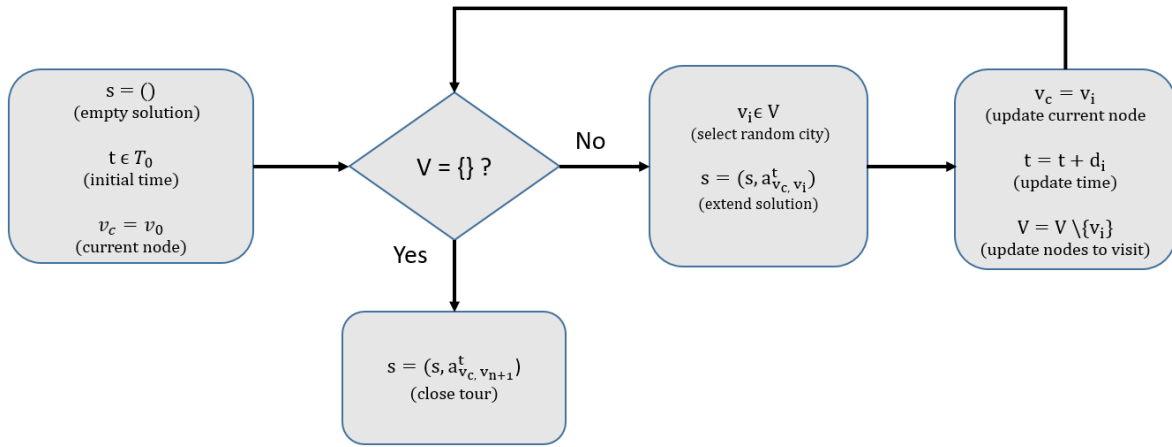
3.2.1   Initial Response

Consider a user request which defines a start period $T_0 = [T_{0i}, T_{0f}]$, a start city, v_0, a return city v_{n+1}, and a list of cities to be visited, V, where upon visiting each city $i \in V$, a waiting periof of d_i days is necessary.

The steps necessary to construct a random solution to this request are illustrated in figure **??**, and can be summarized as follows. Start by setting the initial time, $t \in T_0$, and the current node, $v_c = v_0$. If there are no nodes to visit, that is, if $V = \{\}$, a solution if defined by the flight $a^t_{v_0,v_{n+1}}$. Otherwise, select a random city, $v_i \in V$, and extend the solution with the flight $a^t_{v_c,v_i}$. Follow this with the update of $t = t + d_i$, remove the selected node v_i from V, and set it as the current node. Repeat this process untill all nodes are visited, and conclude by closing the tour.

The described process may start immediatly after receiving a request, because it does not require the information relative to the arcs of the problem. Instead, it generates a completly random solution, composed of several flights, whose information is not yet known. Thus, this solution is passed to the Data Management System, which calls a third-party API to request the relevant information. If the number of flights which constitute the solution is lower than 9, this information can be obtained with a single HTTP request to the API. Otherwise, the number of necessary requests is given by, approximately, $N/10 + 1$, where N is the number of necessary flights.

While the random solution generation may start immediatly after receiving a request, the nearest neighbour heuristic relies on a complete cost matrix, and so can not be initiated before collecting all the information necessary. The nearest neighbour algorithm is closely related to the process described in figure **??**, however, instead of selecting a random node $i \in V$, this node is selected according to its cost. At any point of the process, there is a *current* node and time, which can be used to determine all the possible solution components (arcs), where the one with the lowest cost is selected. Furthermore, if the start time of the request is given by a time window, the illustrated process can be repeated for all allowable start dates.

Figure 3.2: Steps to generate a random solution to a user request.

3.2.2 Ant Colony Optimization

Figure **??** presents a block diagram illustrating the Ant Colony Optimization metaheuristic. The algorithm receives a weight matrix which contains the information relative to the arcs of the bipartite graph describing the problem, aswell as other relevant data, as the duration set D, associated to each city i, of the set of nodes to be visited, V. After the initialisation, and untill the termination condition is met, the algorithm performs a cycle, in which every agent constructs a solution to the problem, followed by a pherormone matrix update, to reflect the search experience of each ant.
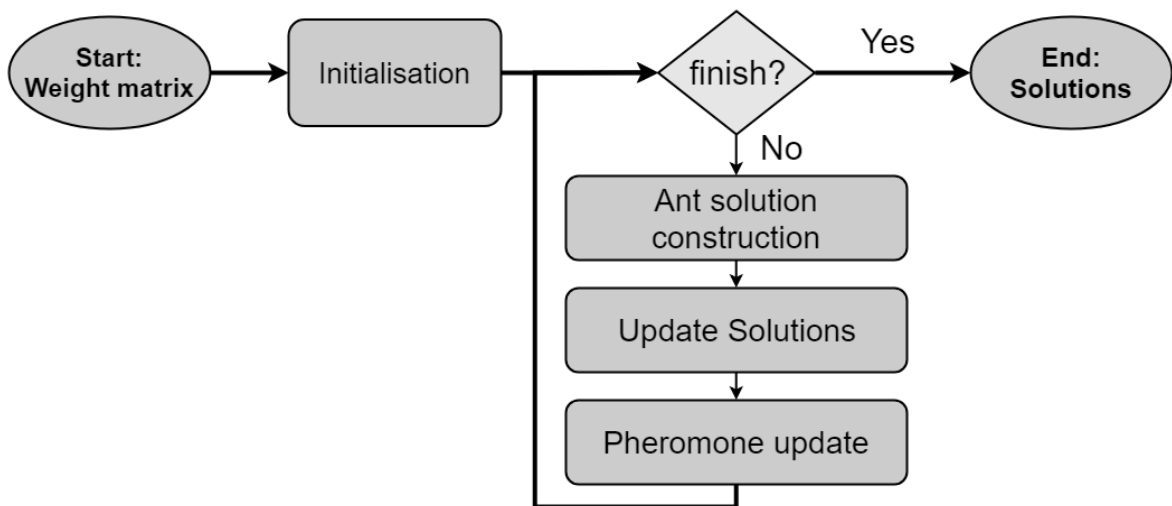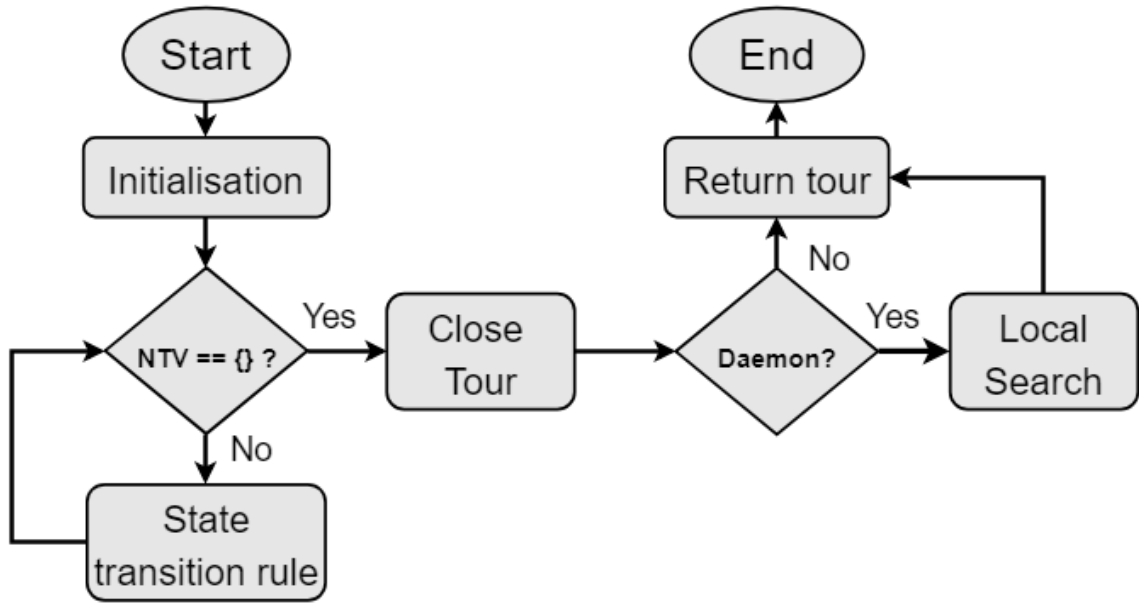

Figure 3.3: Block diagram of the Any Colony Optimization metaheuristic.

The initialisation of the ACO metaheuristic requires the construction of an initial pherormone matrix.

As suggested by the ACO authors [**?**], the value of each entry of the pherormone matrix is given by equation **??**, which is inversely proportional to the cost of the nearest neighbour solution, C^{nn}. The initialisation of the metaheuristic also requires the definition of a variety of algorithm specific parameters, as the number of ants m, the pherormone evaporation rate ρ, the heuristic relative influence β, the pherormone relative influence α, and the exploration rate Q_0.

$$\tau_{ij}^t = \tau_0 = \frac{1}{nC^{nn}} \tag{3.4}$$

The construction process undertaken by each ant, illustrated in figure **??** is as follows. First, the current time is set to a value belonging to the allowable start dates, $t \in T_0$, and the current node is set to the start node v_0. Each ant enters than an iterative cycle untill all nodes belonging to V are visited. At every step of this cycle, an ant chooses a solution component by either *exploiting* or *exploring* the search space. The decision of exploiting or exploring depends on the algorithm parameter Q_0, and a pseudo-random value q, calculated at run time. The selection of the solution component j, which identifies the next city to be visited, is thus given by equation **??**. After the selection of each solution component, it is necessary to update the time, incrementing it by the duration relative to the selected city.

Figure 3.4: Block diagram of the construction procedure undertaken by each ant.

$$j = \begin{cases} exploitation \text{ (eq. \textbf{??})}, & \text{if } q \leq Q_0 \\ exploration \text{ (eq. \textbf{??})}, & \text{otherwise} \end{cases} \tag{3.5}$$

The exploration of the search space utilizes the so called *random-proportional rule*, given by equation

43

??, which determines the next solution component of the ants solution. Note that $J_k(i, t)$ is the set of the solutions components which might be selected and form a valid solution, by an ant which is currently at city i at time t. In its turn, the exploration is given by equation **??**, and $p_a(i, j, t)$ represents the probability of ant a, which is currently at the node i at time t, selecting j as the next node to visit. In the presented equations, η is the inverse of the weight matrix value.

$$argmax_{j \in J_k(i,t)}[\tau(i,j,t)][\eta(i,j,t)]^\beta \qquad (3.6)$$

$$p_a(i,j,t) = \begin{cases} \dfrac{[\tau(i,j,t)][\eta(i,j,t)]^\beta}{\sum_{u \in J_k(i,t)}[\tau(i,u,t)][\eta(i,u,t)]^\beta}, & \text{if } j \in J_k(i,t) \\ 0, & \text{otherwise} \end{cases} \qquad (3.7)$$

Follwing the iterative construction procedure, an uncomplete solution is available. To complete this solution, it is necessary to add an extra solution component, which connects the last visited node, to the return node, v_{n+1}.

Having a complete solution, it is possible to perform *Daemon Actions*, or in other words, try to improve the constructed solution using a local search algorithm. The execution of this step is optional, and is set by a parameter algorithm at the initialization of the procedure. The selection of a local search algorithm is usually problem specific, and for the Traveling Salesman Problem, one of the most used local search procedures is the k-opt exchange. During the development of this work, the 2-opt procedure was utilized.

The mainstream success of the Ant Colony Optimization in the multiple areas in which it was applied, is due to its ability to *guide* the agents to good solutions. While the construction of each solution follows a pseudo-random rule, the overall quality of this solution is used to bias the search towards the search space around this solution. This achieved by the global pheromone update, given by equation **??**.

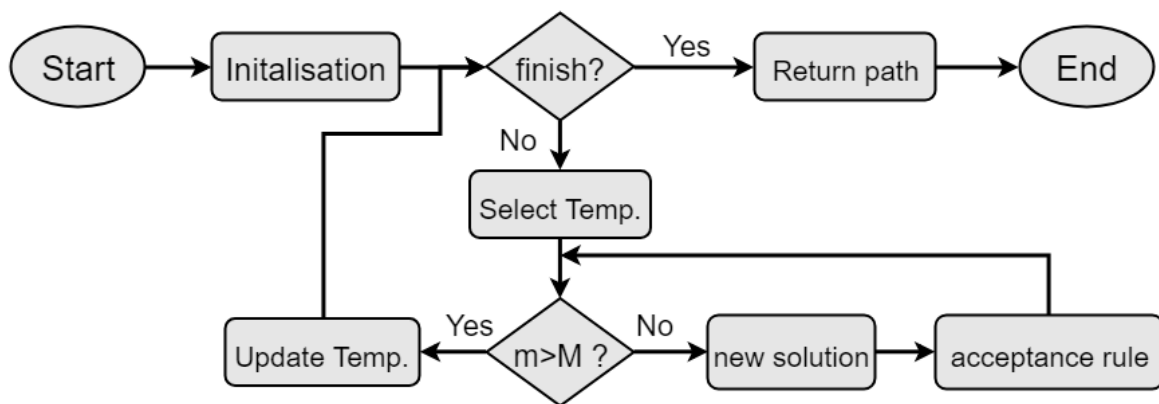$$\tau_{ijt} = (1 - \rho)\tau_{ijt} + \sum_{s \in S_{upd}|c_{ijt} \in s} g(s) \qquad (3.8)$$

The global pheromone update is a proccess in which, at every step of the algorithm, all entries of the pheromone matrix are updated, by the so called pheromone *evaporation* and *depositation*. Pheromone evaporation is a method used to avoid getting stuck in sub-optimal solutions, while pheromone deposition intends to induce the exploration of the search space around good solutions, by incrementing the pheromone values of the solution components belonging to a set S_{upd}. The set S_{upd} is usually a subset of the solutions constructed at the current iterative step, $S_{upd} \subset S_{iter}$. Furthermore, the set S_{upd} includes the best global solution S_{best}, and this work follows the elitist ant rule, which defines that pheromone deposition on the solution components belonging to S_{best} is actually higher than those of the other solutions in S_{upd}. In equation **??**, $g() : S \rightarrow \mathcal{R}^+$ is a function such that $f(s) < f(s') \Rightarrow g(s) \geq g(s')$. It is common set $g()$ to be inversely proportional to the objective function $f()$. Thus, the inversely proportional constant

for g is set to 1 for all solution components in S_{upd}, and to a higher value for S_{best}, 3 in this case, due to the elitist ant rule.

3.2.3 Simmulated Annealing

The Simulated Annealing is a metaheuristic for solving continuous and discrete combinatorial optimization problem, and it is well known for its ability to escape from local minima. Every SA algorithm is composed of a main cycle, and a secondary cycle, which is often refered to as the Markov cycle. During the markov cycle, new solutions are generated, based on the current solution, and, if the new solution is better than the current one, it is accepted, while a worse solution is conditionally accepted, according to the Metropolis acceptance criterion. This key characteristic of accepting worse solutions, also called hill-climbing moves, is what enables the metaheuristic with the possibility of escaping local minima. The probability of accepting a worse solution is influenced by a parameter of the algorithm, called temperature, and is higher for higher values of the temperature. Thus, the metaheuristic is defined in such a way that the initial temperature is high, accepting worse solutions with a high probability, but during the execution of the algorithm the temperature is decreased, and so is the probability of accepting a worse solution. The adjustmement of the temperature is done at each iteration of the main cycle, after running a complete markov cycle. As the algorithm reaches the end condition, the algorithm becomes gradually more gready, and accepts only improvements to the current solution.

Figure **??** presents a simplified block diagram for the Simulated Annealing metaheuristic. This block diagram is valid for the majority of the SA algorithms, because it is constituted by only the main blocks of the algorithm, and does not specify the cooling schedule, neither the solution generation process, nor the acceptance criteria. This three blocks of the diagram are responsible for introducing differentiation between the different SA algorithms.

Figure 3.5: Block diagram of the Simmulated Annealing metaheuristic.

The majority of the SA algorithms operate on a single solution, but there are authors which propose a multi-agent approach [**?**], and mention that the classic SA does not learn from its search history in an inteligent way, as other meta-heuristics, like the ACO, does. Other authors focus on creating a cooling schedule which may benefit the SA algorithm, by enabling a more exhaustive exploration of the search space around more promising temperatures. An example of this is the List-Based SA algorithm proposed in [**?**], which creates a cooling schedule which initially decreases faster than the traditional geometric schedule, escaping faster from the non promising temperatures, and which than decreases slowly, inducing a more exhaustive search around promising temperatures. The implementation of the Simmulated Annealing algorithm developed in this work, follows the multi-agent and list-based chooling schedule, introduced in [**?**] and [**?**].

Following the block diagram illustrated in fig **??**, the algorithm receives a weight matrix describing the problem, and other relevant information, as the duration associated to each city, and possible constraints relating the initial and final node. The algorithm is than initialised, and some parameters must be set. This includes the initial acceptance probability p_0, the maximum length of the temperature list L_max, the main cycle stop criteria and the markov cycle length. The initialisation also requires the construction of an initial temperature list, which is done as follows:

1. Create an empty temperature list L;
2. Create an initial solution x;
3. Create a candidate solution y from x;
4. If $f(y) < f(x)$, swap x and y;
5. Insert $t = \frac{-|f(y)-f(x)|}{p_0}$ into the temperature list;
6. Repeat steps iii) to v) untill the temperature list reaches its maximum length;

The Simmulated Annealing metaheuristic is based on two cycles. The inner cycle, often called the *Markov chain*, is responsible for producing candidate solutions based on the current one. The algorithm may accept or reject this candidate solution, and this depends on the solutions objective function value, aswell as the current state temperature. The Markov cycle usually runs for a fixed number of times, at each main iteration cycle. After completing the Markov cycle, the temperature is decreased, and if the termination condition is not met, the Markov cycle restarts.

The process by which a candidate solution is generated is problem specific. For the Traveling Sales-man Problem, there are several suggestions in the literature, which include, but are not limited to, 2 and 3-opt moves, insertion, reversion and swapping mechanisms. During the development of this work, two stategies were tested. The first follows a simple 2-opt move strategy at each markov cycle, while the second is a greedy approach which selects the best result of the insertion, reversion and swapping functions.

At each step of the markov chain, the algorithm dictates that if a candidade solution is better than

the current one, this solution is accepted, and set as the current one. In its turn, if a candidate solution is worse, it not automatically discarded, but it may be accepted. This set of rules is refered to as the Metropolis acceptance criteria, and is defines in equation **??**. The criteria sets that when confronted with a worse solution, a random number $r : r \in [0, 1[$ is generated, and if r is less than the acceptance probability, $e^{-\frac{f(y)-f(x)}{t}}$, the candidate solution is accepted. This probability is such that, the lower the objective function difference, the higher the probability of accepting the solution. On the contrary, as the state temperature decreases, so does the acceptance probability.

$$p = \begin{cases} 1, & \text{if } f(y) \leq f(x), \\ e^{-\frac{f(y)-f(x)}{t}}, & \text{otherwise} \end{cases} \tag{3.9}$$

The Simulated Annealing convergence theory suggests that this algorithm is capable of reaching the global minima,
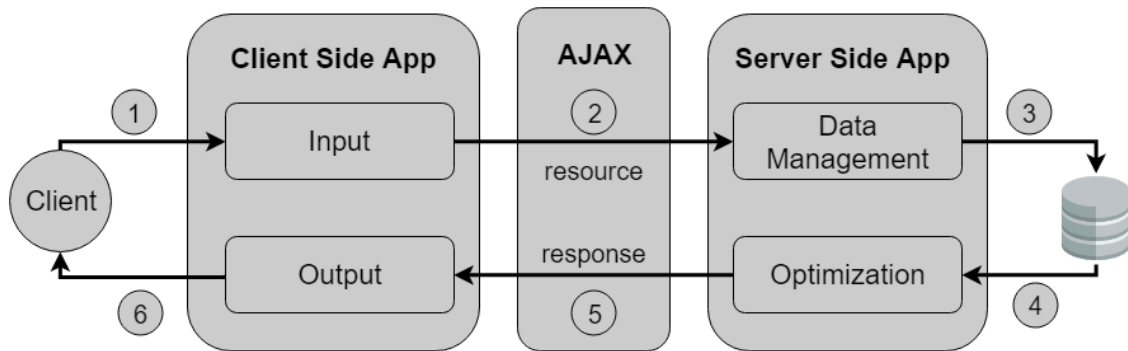
4

System Design

4.1 System Architecture

The proposed system consists of a web application, composed of both a client, and server side application, which are mutual exclusive but dependent upon another. The basic structure and data flow of the proposed application is presented in figure **??**, which will be explained with further detail.

Figure 4.1: Structure and data flow of the proposed application

The client side application (CSA) is designed to interact with the user, enabling the construction of requests (number 1 in the figure), and the presentation of solutions to those requests (number 6). (The details of this design and implementation are covered in section x and y, respectively.)

Due to the nature of the constructed requests, which are Np-hard for unconstrained multi city requests, the optimization process associated to it is computationally heavy. Thus, the response to a request is not processed by the client side application, but rather by an application which runs on the server.

The communication between the client and the server applications is done via Asynchronous Javascript and XML (AJAX), which means that upon submiting the request (number 2 in the figure), the client can continue to interact with the application while the response is being prepared.

The server side application is responsible for producing a response to the user constructed requests. (As covered in section x), the user requests can be transformed into a Traveling Salesman Problem instance. However, upon the construction of the request, only the nodes (cities) of the TSP problem are specified, while the arcs (flights) which connect those nodes, are not. Thus, it is crucial to complete the TSP graph with real fligh data (number 3 in the figure). (This essential step of accessing real flight data may be subject to much discussion, and is covered with more detail in section X). Having access to a complete TSP graph, it is possible to run an optimization algorithm (number 4), which produces the

solution to the user request.

Finally, when the response is ready (number 5), the client application uses this information to update and re-render the page, presenting the result (number 6). (This interaction between both applications is discussed with more detail in section z.)

4.1.1 Client Side Application

The goal of the Client Side Application (CSA) is two interact with the user. Due to the work being developed, the CSA must:

1. enable the construction of user selected requests;
2. interact with the API to request specific resources;
3. display the response to the user requests;

The construction of user selected requests is done in such a way that it can be used to instantiate a Flying Tourist Problem, which enables further optimization. Following this definition, the user interface must enable the collection of:

- the start city, v_0;
- the return city, v_{n+1};
- a list of cities to visit V;
- the durations D associated to each city in V;
- the start time/period T_0 of the trip;

Due to the computational complexity associated to this request, the Client Side Application will not handle the optimization, since this would consume too much of the users device resources. Thus, upon having a complete user request, it is passed to a Server Side Application (SSA) to be processed.

Finally, upon receiving the response of the SSA, the User Interface must be updated so it can display the solutions to the request. A solution to a request is an object which containts at least one set of flights which satisfies the user defined query. However, a solution to a response should contain several valid solutions, so the user might choose the one which is more adequate to his needs. Furthermore, it is important that each of the flights presented contains at least the most relevant information about it. Thus, upon presenting the solution, each flight must include, at least, the following attributes:

- the flight cost;
- the flight duration;
- the date, departure and arrival time;
- the number of layover flights;

At this point, it is well defined what constitutes the user Input/Output, and that the Client Side Application handles only the logic of collecting user requests, and of preseting solutions to these requests. In its turn, the processing of the request is done by a third-party API - the Server Side Application.

Having a clear idea of what the CSA must do, it is possible to talk about its actual design. The views of the User Interface can be grouped into two simple structures, the *Request* view, and the *Response* view. These views enable, respectively, the construction of user requests, and the visualization of the constructed response. There is a third view which would be very interesting to have - a *Map* view. While the Request and Response views are essential to the overall function of the application, the Map view is not. However, a map, which could display the routes of the selected flights, aswell as other relevant information, would certainly contribute to a better and more complete user experience.

It is also important to dicuss the design of a *response* User Interface, that is, a UI which adjust the size of the presented content according to the characteristic of the users device. Today, the browsing of the internet occurs on a multitude of devices, which include mobile, desktop and tablet devices. Each of these classes of products has several ranges of dimensions for the devices included in each category. There are some mobile devices which have 3 inches screen, while others may have 6', and the size of the screen increases considerably as we move towards laptop and desktop screens. Thus, upon rendering a webpage, it is important to know the size of the decive being used, and adjust and resize the displayed content accordingly.

With this in mind, the proposed User Interface should follow the design illustrated in figure **??**. Notice that there are 3 views, the Response, Request, which are essential and thus are always present, and the Map view, which is not, and thus can be discarded on some devices.
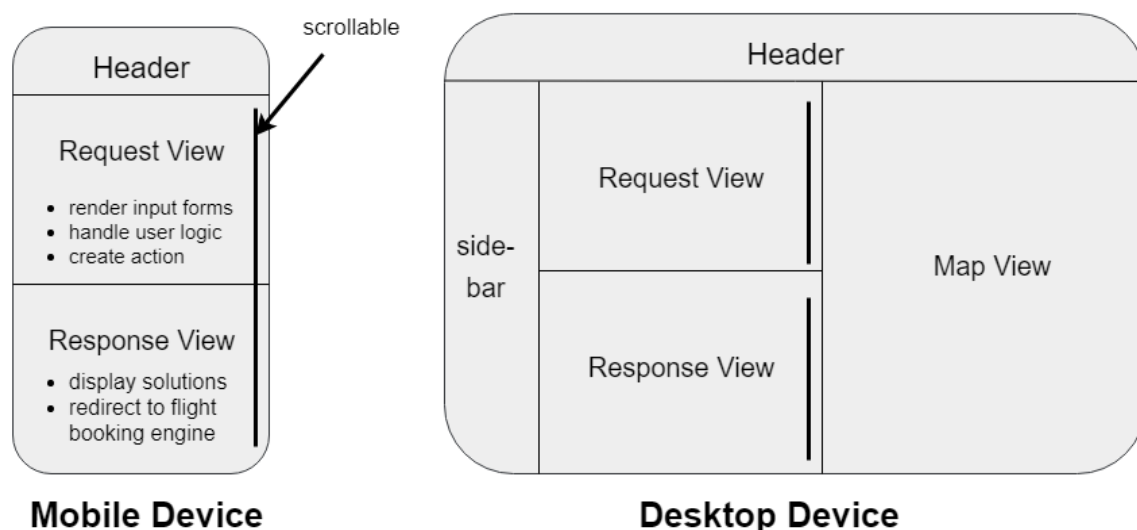
Figure 4.2: Proposed response User Interface for small/medium and large devices. There are 2 essential views and one optional view.

4.1.2 Server Side Application

The Server Side Application (SSA) is the module responsible for creating a solution to a user selected request. In this case, the SSA receives a shallow Flying Tourist Problem, that is, an instance of the problem which does not yet have the set of arcs A associated to the problem. Thus, the SSA has two main goals:

1. create the list of arcs associated to the problem;

2. produce a solution to the user selected request;

The system which is responsible for producing the list of arcs will be denoted as *Data Management System*, while the production of a solution to the user request relies on a system called *Optimization System*.

4.1.2.A Data Management System

Upon receiving a user specified request, the set of nodes, aswell as the start time and durations, are well defined. On the other hand, the set of arcs which connects these nodes is not. For example, a user request may correspond to a single flight between A and B at time t, and upon receiving this request, there is no information available about the flights (arcs) which connect these two cities. In fact, thats exactly what the user is looking for. Thus, the goal of the Data Management System is to collect the necessary information to construct the list of arcs associated to a user request.

It is important to note that an arc connecting two nodes corresponds to a flight between two cities, at a specific date. However, there are multiple flights which fit this discription, and every one of these flights has several attributes which differentiate themselves. For exemply, every flight has a particular cost, flight duration, departure and arrival time, airline, bag limit or even layover flights.

Due to the vaste attributes which define every flight, it is impossible to known which particular flight is the most adequate to a specific user, because users often have different selection criteria. Due to this, upon the construction of the multipartite graph, it makes sense to have a list of possible flights, for every arc in A, instead of just one. This enables the selection of a specific flight according to the objective function being minimized. For example, if the goal is to minimize the total fligh cost, it makes sense to select those flights which present the lowest cost, disregarding other attributes of the flights as, for example, the flights duration. This means that when talking about an arc connecting two nodes, it makes more sense to talk about a *family* of arcs, which share key characteristics, as the origin, destination and date, but which may vary regarding the other attributes, as flight cost and duration.

Following the definition of the Flying Tourist Problem, and the notation proposed by J.C.Picard, a request may be defined as a multipartite graph, divided into k layers, and where every node in a layer

is fully connected to every other node in the subsequent layer. The total number of layers, which correspond to time span between the earliest date in which the trip might start, and latest date in which it might finish, is given by equation **??**. The arcs which connect those nodes are divided into three groups: *initial*, *transition* and *final* arcs. Figure **??** illustrates the time periods associated to each arc type, according to the Flying Tourist Problem definition, with a multiple time start period, and a non empty list of cities to be visited.

Figure 4.3: Illustration of the distribution in time of the initial, final and transition arcs.

$$k = T_{fb} - T_{0a} + 1 = t_s + summ(D) + 1 \qquad (4.1)$$

The initial arcs correspond to those who might initiate the trip, and thus must start at city v_0 at a time $t \in T_0 = [T_{0m}, T_{0M}]$ and visit every city in $V \cup v_{n+1}$. There are a total of $k_1 = T_{0M} - T_{0m} + 1 = t_s + 1$ layers for the initial arcs.

In its turn, the transition arcs are those which connect the N cities of the list of cities to be visited V. The earliest transition arc occurs at a time no sooner than $t_1 = T_{0a} + min(D)$, where $min(D)$ corresponds to the lowest entry of the set of durations. That is, if the trip starts at time T_{0a}, which is the earliest start date, and if the minimum duration associatedto a city is $min(D)$, than the second flight might occur at a time no sooner than $t_1 = T_{0a} + min(D)$. Following a similar approach, the latest transition flight can occur no latter than $t_2 = T_{0b} + summ(D) - min(D)$, where $summ(D)$ corresponds to the summ of all entries belonging to D. Thus, there are a total of $k_2 = t_2 - t_1 + 1$ transition layers.

On the other hand, the final arcs are those which might conclude the trip, and thus must return to city v_{n+1}, from a city in $V \cup \{v_0\}$. The final layer extends from T_{fa} to T_{fb}, where $T_{fa} = T_{0a} + summ(D)$ and $T_{fb} = T_{0b} + summ(D)$. There are a total of $k_3 = T_{fb} - T_{fa} + 1 = t_s + 1$ final layers. Comparing to the number of initial layers, one can conclude that there are as many initial layers, as there are final layers.

In conclusion, the goal of the Data Management System is to enable the construction of the lists of

flights which are necessary to proccess some user request. By now, it should be clear which particular set of flights are necessary. However, nothing was yet said about how to actually obtain the required information.

To have access to real flight data, the most simplest and efficient way is to find a publicly available flight data API. There are several choices regarding to this, and it was a subject of a lot of research. In summary, the available API's are classified as *free*, *limited* and *unacessible*. A free API is one which charges no charge for any query, nor limits the number of daily queries. On the other hand, a limited API is one which sets an upper bound on the number of daily available queries, and charges a fee after this limited. In its turn, an unacessible API is an API which is available only for commercial solutions, and which was unavailable to be used in a research context. In this work, only free API's were tested and used. The communication to an API relies on simple HTTP protocol, where the requested resource is specified according to a URL syntex defined by the source, and whose response is usually in the format of JSON, corresponding to a list of flights, with a lot of details about the attributes of each flight. Figure **??** illustrates the response of a particular API (Kiwi) to a specific query. In this image it can be seen that there are a total of 134 flights satisfying the query (data [113]), and that each flight has multiple attributes, as the price, duration, country etc.
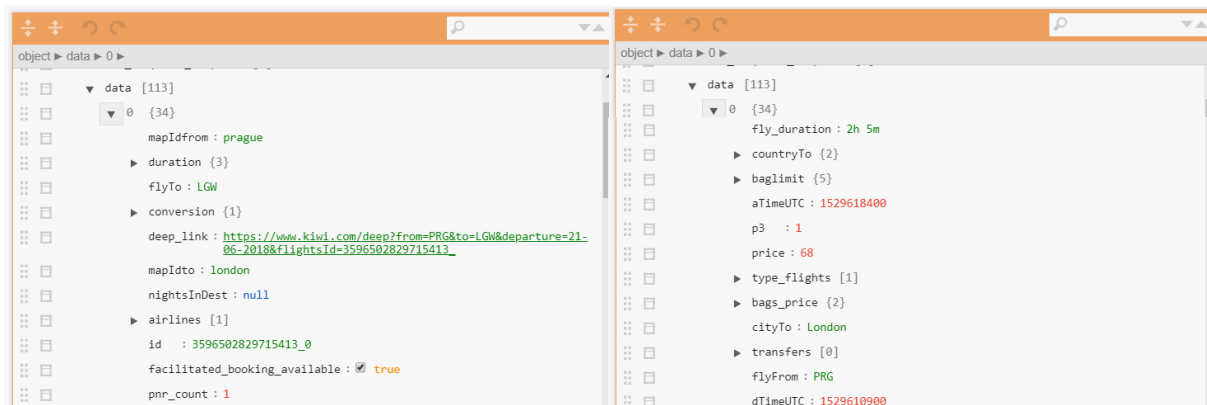

Figure 4.4: Response of Kiwi's flight API to a particular query. It can be seen that multiple flights are returned, each with its own attributes.

One of the most important conclusions of testing different API's, is that different API's provide very different results. These results are presented in annex **anexo preços aqui!**, and show that for the same query, flight prices may vary up to 103%, depending on the source used.

To tackle this problem, a secondary approach was taken. Instead of using public API's to access real flight data, it is possible to create a program known as *web scraper*, which visits a specific website in order to retrieve the information from it. This enables the consultation of flight sources, which otherwise could not be acessed. While webscraping has the advantage of accessing content that otherwise could not be, it has many disadvantages. Comparing to API's, web scraping is a lot slower, and much more

difficult to actually implemenet. Figure **??** illustrates the process of inherint to web scraping. On the left is an image of a website query and results, while on the right is the corresponding (simplified) source code. The annotations provided identify the different flight attributes, and the locations in which they occur in the HTML code.
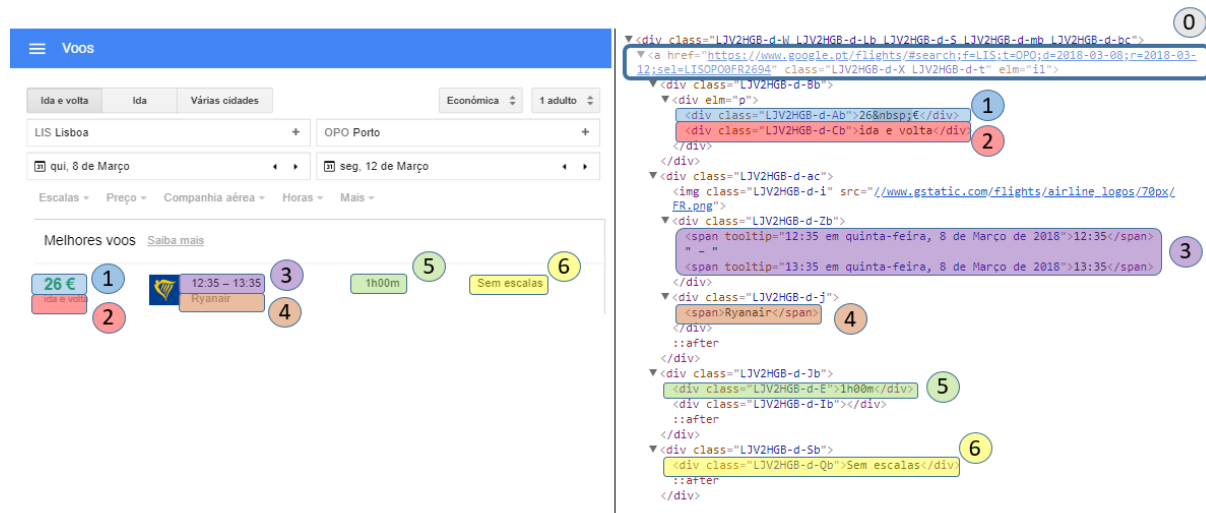
Figure 4.5: Website User Interface displaying a flight query, and corresponding source code. A webscraper can use a parsing method on the source code in order to access and collect this data.

4.1.2.B Optimization System

The goal of the optimization system is to produce a solution to a user defined request. When a user defines a request, as shown before, the arcs which connect those nodes are not defined, and thus, no solution can be produced. After collecting the information regarding these arcs, it is possible to run an optimization algorithm which produces the best set of flights for the specific request, according to some objective function. However, depending on the specific request, the time necessary to collect the required flights, or the time necessary to produce a high quality solution may be very high, due to high number of necessary flights, aswell as the computational complexity associated to this problem.

Because of the latency inherint to complex requests, the optimization system will try to produce a random solution, whose overall quality is unknown, in a very short time. This can usually be done using a single HTTP request, because API's allow the aggregation of up to 9 flight queries, and it enables the presentation of a very fast response, reducing the latency felt by the user. Furthermore, this type of random search corresponds to the currently available option for multicity flight searches. That is, every website which does multi city search, require the definition of a specific route and specific dates, and it is this approach that will be taken to produce a first random solution.

Of course, the random solution produced is probably not the best, because it is only one of the more

than $N!$ possible solutions for a N city multi trip request. Thus, after producing the initial response, the optimization system will utilize more specialized algorithms to try to produce better solutions. These algorithms usually require a complete cost matrix, and thus this process can only be undertaken after the Data Management System completed the construction of the family of necessary arcs.

Having the complete list of arcs, it is possible to run more complex optimization algorithms, which are usually capable of producing better solutions, but which might take some time to do so. A first simple heuristic algorithm that can be produced is the Nearest Neighbour, which is probably capable of improving the quality of the previsouly defined random solution, in a very short time. On the other hand, two meta-heuristic algorithms, the Ant Colony Optimization and the Simmulated Annealing, will also be tested. The objectives is to improve the quality of the solution as much as possible, in a reasonable time.

This process of solving the request using multiple optimization algorithm is illustrated in figure **??**. Note that the *Random* solution can be presented in a very short time, because it is not necessary to construct a complete multipartite graph, but simply a single random solution, which can almost immediatly be presented to the user. On the other hand, the *N.N.* and *Metaheuristic* solutions can only be produced after having a complete multipartite graph. Thus, it is expected that producing these solutions may take much more time compared to the random.
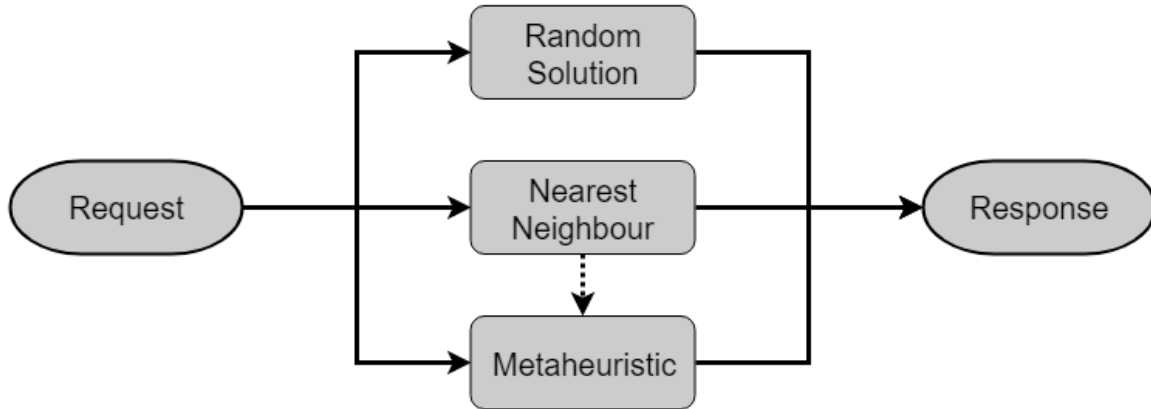
Figure 4.6: Simplified illustration of the optimization system, which utilizes different algorithms to produce a solution to a user defined request.

Having an optimization system which employes multiple, very different, optimization algorithms, enables the comparison of both the time, and the quality of the proposed solutions. We known apriori that the random solution is fast, and that its quality is probabily bad. On the other hand, we expect the more specific algorithms to be slower, but present better solutions. The goal than to verify if the solution is actually better, how significant the improvement is, and how much time was necessary to produce such an improvement.

4.2 Design Choices

This section wishes to clarify some of the design choices considered, and the reasons that led to particular decisions.

4.2.1 Problem Formulation

The formulation of the problem relies heavily on the goal of the overall system. The start point to this work was the development of an flight search application, which could solve the multicity trip routing problem, aswell as any single and round flights, presenting the best set of flights, that is, those which would minimize the total cost of the trip.

To develop rich user experience, that is, a search which is simple and declarative, and covers a broad number of situations, it is necessary to think like a user. No matter what trip the user is searching for, there are always some questions which are the same, and can be grouped into 4 classes: the where, when, what, and how. That is, *where* do we start and finish, *when* do we start, *what* cities will are to be visited, while the *how* may represint a variety of things which further characterize the trip, as the duration, possible time windows and/or routing constraints.

In most cases, users have a specific start and returning point, which might or not be the same. In its turn, the schedule for the trip is usually much more broader. That is, instead of having a single start date, in some cases user consider a time span of a few days, weeks, or even months. Thus, it was decided that instead of forcing a single start date, a time period would be considered. The same type of reasoning can be applied to the duration of each city. While some users may have a well defined number of days, say 3, others might be more flexible and consider staying in the city for 3 to 5 days for example. This flexibility upon the duration was considered during the development of the formulation, but not introduced in the overall system, due the extra complexity it would introduce, and also because it would be difficult to effectively test such a system. The same can be said about the time windows associated to each city. While it makes sense to have a broader definition which contemplates this possibility, it is more convenient to have a simpler but efficient formulation, which covers only the most important and global search criteria.

4.2.2 System Goal

While the goal of the system was always the development of a flight search engine which would always return the cheapest set of flights, upon having a complete system it was noticable that in some cases, these set of cheapest flights could be dismissed by the user, for having some other bad attributes. That is, a flight could be very cheap, but at the expense of having a very high flight duration.

To tackle this problem, we propose multiple objective functions which cover different goals. On one hand, we can define an objective function for the minimization of the total cost, and upon constructing the TSP problem, we consider the flight cost as the cost matrix entries. On the other hand, a similar approach can be done to minimize the total flight duration. Intead of constructing a cost matrix with the flight cost, this is done with the flight duration. Finally, the same approach is considered for the minimization of the *entropy* of the trip. We borrow this concept, entropy, to describe the weighted summ of the flight cost and duration. Selecting the set of flights which have the lowest entropy leads to the increase in the *efficiency* of the proposed solution.

Having defined three objective functions, any user request will be responded to with a list of three different set of flights, which correspond to the cheapest, fastests, and most efficient solutions found.

4.2.3   Data Collection and Storage

As noted several times, solving a user defined request requires the access to real flight data. During the development of this work, two different forms of doing this where tested: API access and web scraping. Having access to a flight API is the standart choice, it is fast, customazible and efficient. However, the number of available free API's is very low, and overall, the results presented by each may vary considerably. More than this, the problem is that some API's limit the content they offer, not including, for example, low cost airlines. This problem was faced specifically with Google Flights Api. To tackle this problem, and to actually find the cheapest flights, an alternative was considered, web scraping.

Web scraping, or screen scraping, is the act of visiting a page and copying its content. This enables the parsing of the page, and the exctraction of usefull information, which can than be stored into a database, or other form of storage, in order to be accessed. Thus, web scraping can be used to collect information about the necessary flights, however, as it was later verified, it is a very complex and specific process. Furthermore, it is also considerable slow, because most websites use javascript to produce their information, and this usually requires some access to their database, which is in general very slow.

Considering the advantages and disadvantages of each system, it was decided that the developed system would use public API's to access flight data, because it is much faster and more efficient.

User requests are self contained, that is, the data necessary to process them is gathered, used, and than discarded. It would be possible to store each collected flight into a database, for later consultation. However, this data would have an extremely short *expiration date*, or, in other words, the cost of a particular flight might vary at any moment. Thus, it can not be proved that the previously stored data is still accurate and up to date. Because of this, it was opted to not implement a database. This choice might be adapted at any point, and should be considered if the system actually scales to a point in which there are multiple equal queries each day.

4.2.4 Optimization System

Regarding the optimization system, the goal was always the development of a time efficient algorithm. Producing a complete multi city search is a complex problem, which can take a considerable time. However, waiting for a considerable time is not an option for most users. Thus, it is important to produce a solution as soon as possible, and try to improve the solution in the subsequent moments. This lead to the system proposed in figure **??**.

4.2.5 User Interface

Upon the construction of the client side application, which handles the logic of collecting user requests, aswell as the rendering of the user interface, it was important to choose a design which would fit both mobile and desktop users. There are several market study reports which show that the mobile market segment is, and will continue to, experience a growth in its size, textbfcite. Thus, the design of the user interface had to be prepared for the multitude of user devices, or in other words, it had to be *responsive*, textbfcite.

5

System Implementation

5.1 System Architecture

The developed applications are divided into two distinct groups: the Client Side Application (CSA), or the User Interface, and the Server Side Application , or the flight data/optimization API. Both applications run independently from one another, although the CSA is dependent of the API, and must communicate with it as to obtain solutions to the user requests. Figure **??** introduces the architecture of the system, aswell as the technologies on which the system relies, which will be adressed in the rest of this section.
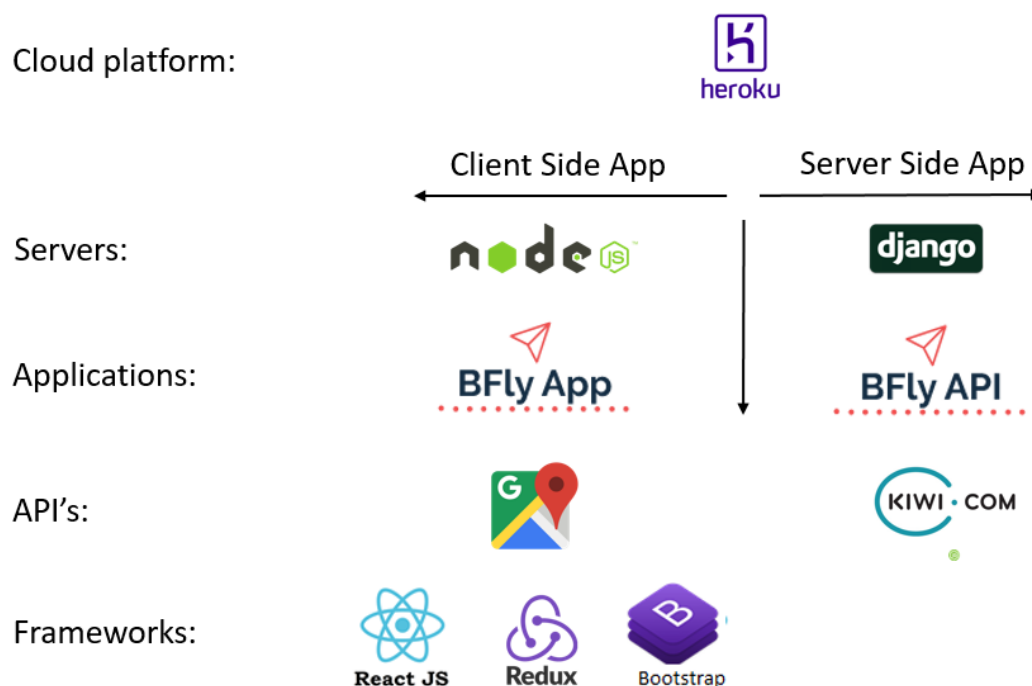
Figure 5.1: Architecture of the application.

The developed applications run on *Heroku*, a Cloud Platform as a service, which, upon request, creates two separate servers, one for each application. The Client Side Application runs a *Node.Js* server, which creates a bundle file, that is served to the user, and contains all the necessary information to render the UI, handle the user input logic, and interact with the API. In its turn, the API runs on *Django*, which interacts with the webserver to read the user request, and may execute particular intstructions according to the selected route.

5.1.1 Underlying technologies

5.1.1.A Heroku

Heroku is a Cloud Platform as a Service (PaaS), which allows applications to be built, deployed, monitored and scaled, in an easy and fast way, by bypassing implementation details specific to infrastructure

and software, as the hardware and the servers [**?**]. In terms of services, Heroku competes directly with other cloud platform services, which include, but are not limited to, Google's App Engine [**?**] and Amazon's Web Services [**?**]. A succinct comparison of the advantages and inconveniences amongst these different services can be consulted in [**?**], while a more comprehensive (although somehow outdated) overview is available at [**?**].

Heroku provides a detailed explanation of its services and how to use them in order to deploy and manage an application, available at [**?**]. Following this overview, an Heroku application can be defined as a *slug* which runs on *dynos*. Clarifying these terms, a *slug* is a bundle of source code, any required dependencies, the programming language runtime, and any compiled output of the build system, everyting ready for execution. In its turn, a *dyno* is an isolated, virtualized Unix container, that provides the necessary environment to run the application.

A slug is typically the result of the deployment and bundling of an application. The Heroku client has a particular application defined by its source code, and in order to run it on Heroku, he must specify the app's programming language, its dependencies, and a *Procfile* - a file which lists the commands to be executed. These applications can be deployed to Heroku via *git*, or using Heroku's own API. Upon deployment, Heroku builds the app, producing the slug, ready to be executed. These are the fundatamental steps to deploy and run an Heroku app, but it is much more customizable than this. For example, it is possible to specify add-ons, as databases, and e-mail services.

The produced slug runs on the dynos, and Heroku offers the possibility to run applications in a completly free manner, by providing a single free dyno. This free option has one inconvenience, which is that upon 30 minutes of inactivity, the app "sleeps". This means that the next connection to an inactive and sleeping application will require some extra time. It is possible to overcome this by utilizing a paid dyno. Furthermore, as the application grows and the number of users increase, a single dyno may not be enough to handle all user connections. Thus, it is possible to *scale* the app, by increasing the number of total dynos utilized.

During the development of this work, Heroku is used as the host service for both the client and the server side applications. Both of these applications run on a single free dyno, and do not utilize any database or other add-ons.

5.1.1.B Node.js

Node.js [**?**] is an open-source and cross-platform JavScript run-time environment, which executes JavaScript code on the server-side [**?**]. Node enables the development of fast and scalable web servers using Javascript only, by utilizing an asynchronous event driven architecture. Thus, Node.js enables the unification of web development around a single programming language, which may be used both on the client and the server side. Node.js includes the possibility to use third-party packages, or modules,

which can be used for the management of networking, file system I/O, data streams, cryptography and much more. These modules are managed and acessible through the Node Package Manager, *npm*.

Upon the creation of the first browsers, JavaScript was utilized as a scripting language, that was used to modify, at run time, the Document Object Model (DOM) of a webpage, enabling the creation of the first dynamic webpages. Today, due to Node.js, this scripting language is not restricted to the browser, and can be used in the server to create dynamic web pages, even before the page is served to the client [**?**].

Node.js was created by Ryan Dahl in 2009, and upon its presentation in the European Javascript Conference, it utilizes an asynchronous event loop, a low-level Input/Output API, and Google's V8 Javascript engine. This last technology, the Google's V8 engine, is a fundatamental part of the Node.js stack, because it allows the compilation of javascript source code into native machine code, instead of interpreting in real time, as occurs in the browser.

In the development of this work, Node.js is utilized as a way of creating a webserver and serving the User Interface. Upon a client connection request, the node webserver responds with a javascript bundle file, which containts all the necessary information for the browser to render, manage and update the UI upon interaction with the user.

5.1.1.C Django

Django [**?**] is a python Web framework, designed for rapid development and a clean and pragmatic design. It is an oper-source and high-level framework, with dedicated modules for security and performance, and for the execution of repetitive task, such as logging, request management, database access, cookies and session management [**?**]. Django is constructed with a Model View Controller architecture in mind, or following the Django nomenclature, a Model View Template (MVT) architecture.

Following the MTV design [**?**], upon receiving a request, the webserver first interprets the request based on the url, using regular expressions. Upon match, the request is passed to a specific *View*. A view is the heart of the application, specifying the functions to be executed, and returning a response to the request. If the application is database driven, as most Django applications are, each view may dispatch some *model* construction process, which queries the database and constructs the necessary data. Finally, a *template* is populated with the data model, and a response may be returned.

During the development of this application, the server side application is built using Django. Despite this, most of the functionality that Django presents is not utilized, because it is not completly necessary. Since the developed application does not directly utilize a Database, the *model* part of the MTV architecture may be skipped. Furthermore, the Django's template engine may also be completly ignored, because the response to a request is a simple python dictionary/JSON object. Thus, the client side application consists in a simple API which interprets requests according to the specified url, invoking the necessary functions, and produces a response to each request.

5.1.1.D   React and Redux

React [**?**] is a JavaScript library for building and rendering User Interfaces. React is based on some core principles which dictate the architecture of every UI built using this library, which include the concepts of *components*, *state*, *JSX* and *Virtual-DOM* [**?**].

Any User Interface built using React is the result of the *render* method invoked by each of its *components*. The content that is actually presented by a component is the result of its source code, written in *JSX*, which is a combination of JavaScript and HTML to, dynamically, produce valid HTML elements. Usually, each of these elements depend on the *state* of the application, at each moment. Thus, every time the application state changes, the User Interface is updated. Since re-rendering the entire DOM is very expensive, React uses the so called Virtual-DOM to re-render only those components which are necessary. At every state change, the DOM is compared to the virtual-DOM, and only a branch of the DOM is updated.

As applications grow bigger, it becomes increasingly difficult to manage the state of each component, as there are natural dependencies amongst components. Currently, the natural response to bypass this difficulty is to use a JavaScript library called *Redux* [**?**], even though there are other solutions developed for this effect. Redux is a state container, which is often said to be predictable, because it forces the usage of immutable data types. This means that every state change is easily identifiable, which is very convenient for React, due to its Virtual-DOM nature.

Furthermore, Redux is said to be a predictable state container because every state change requires the dispatching of a particular *action*, which triggers a function that manipulates the state, called *reducer*. Thus, Redux applications may implement a "history" of the state, because there is information available about every action that were executed to reach the current state.

During the development of this application, React and Redux were used together to build the client side application. This means that React is responsible for rendering the User Interface, while Redux manages the state, updating it on user actions, and interacts with the client side application using simple HTTP protocol.

5.2   Client Side Application

The objective of the Client Side Application is to render a User Interface capable of interacting with a user, enabling the collection of flight requests, and presentation of usefull solutions. In order to be able to achieve this goals, the User Interface is built using the *React*, the application state is managed through *Redux*, and the responses for user requests are processed by a third-party API, the Server Side application.

Any application built using React and Redux evolves around a concept called *State*. The state of

the application is stored inside the Redux Store, and it containts all the relevant data regarding the current available information, which, in this case, is summarized by the data collected from the user, and any data collected from third-party API's. The complete state cycle of the developed application is summarized in figure **??**.
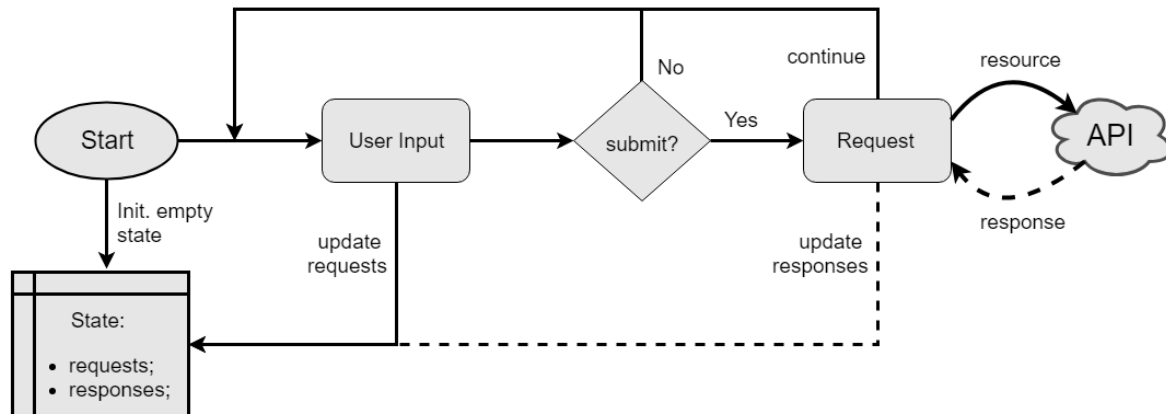

Figure 5.2: Block diagram of the state cycle of the Client Side Application.

Although the user contributes indirectly to a state update, by, for example, filling or updating a certain form, the state can only be modified by a class of functions called *Reducers*. Reducers are pure functions, which means that they always produce the same result for the same input. Each time an action is dispatched, a reducer catches that action, and updates the state of the application. This is usefull for many reasons. To give an example, consider the act of filling a form with the information regarding an airport. The user introduces may introduce a city name, but the reducer may complment this with more meaningfull data as, f.e., the airport ICAO code.

The views of the User Interface are created by the React library, based on the current state of the application and, each time the state is updated, so is the UI. In order to have all views updates with the current state, React proposes a top down hierarchy, divided into *Containers* and *Components*. Containers are on the top of the hierarchy, and are directly connected to the redux store, receiving the current state of the application. These containers usually do not handle presentation logic themselves, but invoke components which do. These components may receive parts of the state, *props*, which may be relevant for rendering purposes, as, for example, the already introduced input of a user form, or the response to a user request.

Although a React application is always up to date with the current state, it does not re-render the entire page every time the state is updated. Instead, it compares the current DOM structure to a virtual DOM introduced by React, and indentifies which components must be updated. This enables a fast and effective update of the User Interface.

Figure **??** introduces the complete architecture of a React/Redux application, including some of the concepts previously discussed. This figure illustrates the top down hierarchy, by having a single component in the top of the hierarchy, which instantiates the store, and renders the complete application, by invoking the containers which are connected to the relevant presentational components. This figure also illustrates the interaction with the store, and with third-party API's. Finnally, it is important to note that a browser can only render HTML, and the entirety of the application is built using a javascript library. Thus, it is necessary to inject the created javascript with all the business and presentation logic, into the html file served to the user. This is usually achieved by using a package compiler called *Webpack*, **cite**.
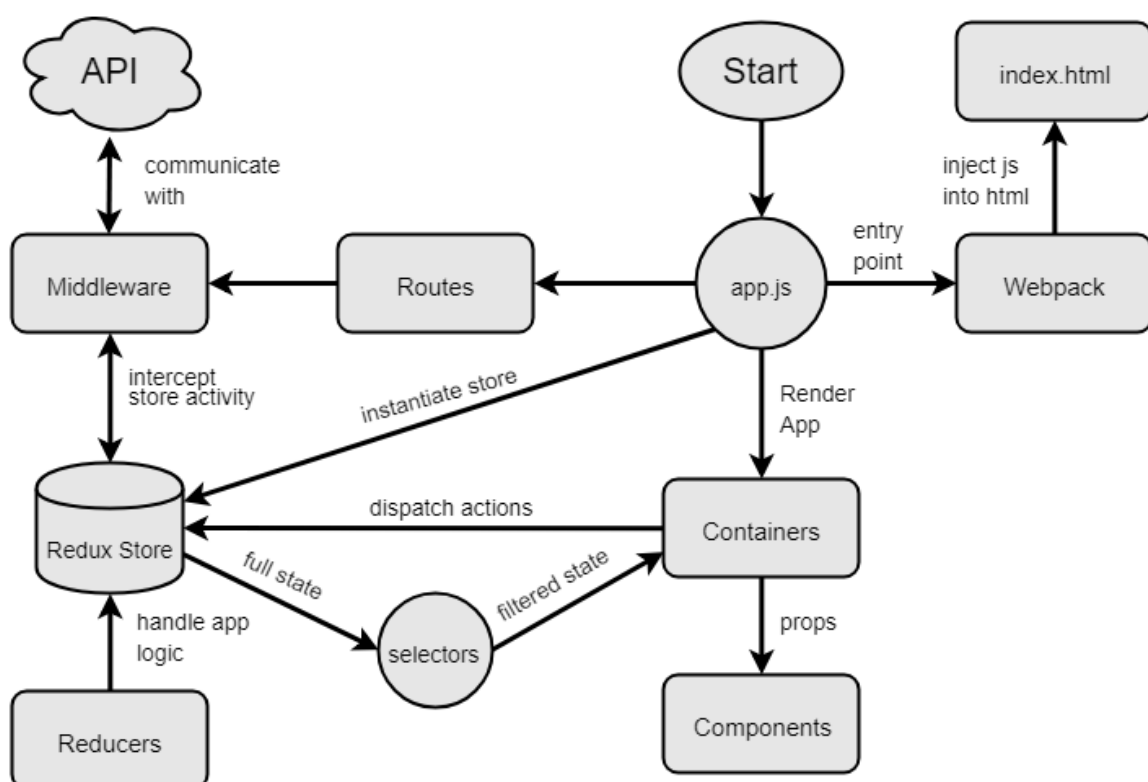
Figure 5.3: Building blocks of an application built using React And Redux

5.2.1 User Input

The Client Side Application intends to solve use requests, and thus, it is necessary to interacting with the user in order to build a complete, valid and accurate request. There are 3 types of requests that will be processed by the application, the single flight, round flight, and multicity trip. A user must initially select the intended type of search, and follow this with the completion of a series of forms which collect the relevant input.

Table 5.1: Parallelism between User Input, Actions and Reducers. To each user defined input corresponds an action, declaring the intent of changing the state with some specific data, and a reducer, which actually modifies the state.

User input	Action	Reducer
origin	actOrigin(origin)	setOrigin(state, action)
destination	actDest(dest, index)	setDest(state, action)
duration	actDur(dur, index)	setDest(state, action)
start date	actDate(date, index)	setDest(state, action)
submit	actRequest(request)	setResponse(state, action)

Any request type requires at least an origin and a destination, aswell as a start date. There are other search attributes which might be relevant, according to the selected trip type, as is the duration associated to each city to be visited.

To every user input is associated an action and a reducer. An action, which corresponds to a plain javascript object, described by an action type, and other relevant content, as the user input, is dispatched each time the user completes or updates an input form. In its turn, a reducer, which is a pure function, is responsible for translating the input, processing it, and updating the application state accordingly.

The implemented application also forces the user to submit the request, by clicking on a button which dispatches an action. During the development of the application, the possibility of removing this button, and automatically dispatching a request, was considered, but rejected, because of the impossiblity of knowing if a certain request is complete or not. For example, given a single flight request, knowing if the request is complete is simply a matter of verifying if an origin, destination and departute date exists. However, for multicity trips, because there is not apriori information regarding the number of cities, the same can not be said. Thus, it must be an user defined action that declares the intention of submitting a request.

Table **??** introduces the relevant data that may be collected from the user. It also defines the action that is dispatched each time the input is updated, and the reducer which is responsible for updating the state of the application. Note that the last user input, the *submit*, is processed by an assynchronous action. This means that upon submiting the request, a third-party API is called, and only upon receiving a response may the reducer be invoked, storing the received data, enabling it to be displayed back to the user.

### 5.2.2 Communication with API

Upon the construction of a complete request, the Client Side Application must call a third-party API to process the request. Since there is no apriori information regarding the ammount of time necessary to process the request, the implementation of this crucial step should be asynchronous. This means that

users can continue to interact with the application, while the request is being processed.

The businnes logic of the developed application is managed by Redux, but Redux by itself is not able to create asynchronous behaviour. In order to do so, it is necessary to use two secondary libraries: redux-thunk, and superagent. *Redux-thunk* is a store enhancer, or a middleware, providing additional functionalities to the store, and *superagent* is a library for asynchronous javascript and XML (AJAX) requests.

While Redux defines that an action must return a pure javascript object, Redux-thunk overwrites this behaviour, and lets an action return a function instead. Thus, an asynchronous action is an action which upon being completed, invokes a function, which dispatches a secondary action. This means that an asynchronous action may update the state two times: the first when the action is initially dispatched, and a second time when the request is complete. In order to inform the user that the request is being processed, the primary action should update the state with some waiting information, while the secondary action interacts with the API.

The details of the asynchronous action previously explained are illustrated in figure **??**. Note that an asynchronous request starts with an user defined action, dispatched from inside a component, and updates the Redux store at two particular moments in time: the first immediatly, and the second after receiving the asynchronous response.
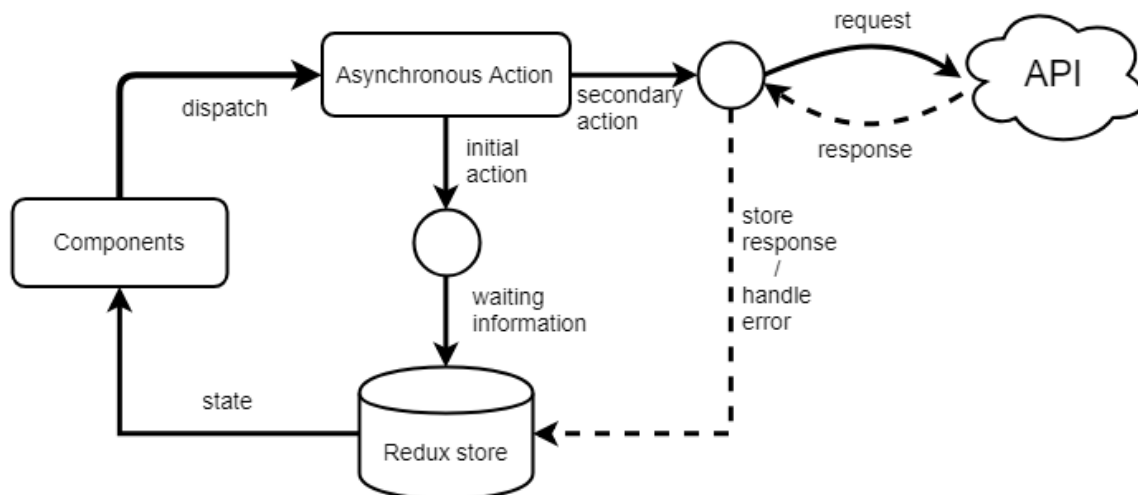
Figure 5.4: Asynchronous javascript request.

The actual request to the third-party API corresponds to a simple HTTP request to a particular URL, according to the API protocol defined in subsection **??**. The URL extenstion, called the Uniform Resource Identifier (URI), enables the specification of the particular resource under request. Thus, the URI may be implemented as a way of specifying the attributes which characterize the user selected

resource. It follows thats a particular resource is identified by a collection of (key, value) pairs, which identify the resource attribute and its user defined value. There are a total of 7 keys which may be used to construct a request: flyFrom, returnTo, flyTo, startDate, endDate, duration and tripType.

5.2.3 User Interface

This subsection specifies the implementation details of the User Interface, and should be considered together with figure **??**, which illustrates the design proposed for the application. The User Interface consists in a single page application, divided into three main views: the *Request*, *Response* and *Map view*. Due to the implementation using React, every view is managed by a container, which reads the state from the store, calls the rendering of presentational components, and may dispatch actions on user input or other events.

The User Interface is designed to be mobile friendly, by being responsive to the device size. This is achieved using the Bootstrap grid system, a website design paradigm in which the user screen is divided into 12 columns, and each block of the user interface may specify a variable number of columns, depending on the screen size. The results of this implementation is illustrated in figure **insert mobile vs desktop figure here** where the mobile and desktop versions are put side a side for comparison.

Figure **??** is a screenshot of the current version of developed application.

5.3 Server Side Application

The Server Side Application is the system responsible for producing a solution to a user request, which corresponds to the specification of a resource, as described in section **??**. Producing a solution to the user request involves the communcation with third party API's, to obtain the necessary flight data, which is handled by the Data Management System, detailed in section **??**. The actual production of a solution is managed by the Optimization System, described in more detail in section **??**. The architecture and implementation details of the SSA are presented in section **??**, and illustrated in figure **??**.

5.3.1 API architecture

The goal of the Server Side Application (SSA) is to proccess user defined requests, and constructing a solution to it. There are two necessary steps in order to actually produce a response to a request: i) data collection, and ii) optimization. Both of these steps are managed by the SSA, respectively in the Data Management System (DMS), and in the Optimization System (OS). The data collection procedure will be explained with more detail in subsection **??**, while the optimization procedure is much more complex, and thus requires a complete section for itself.

Figure 5.5: Print screen of the User Interface, as is, displaying the request of a single flight between Lisbon and Beijing, over the span of a month. Three results are presented, the cheapest (313€/55:10h), fastest (667€/11:35h), and most efficient (368€/11:35h) flight.

The Server Side Application is built using Python, and a Python framework called Django. While the DMS and OS are built with pure Python, it is necessary to wrap these functions around a service which is listening to requests to the application, and this is managed by Django. Django enables the creation of a set of routes relative to the application, and each of these routes is connected to a particular set of instructions. It is also possible to define a pattern each route, enabling the setting of user selected input.

Upon receiving a request, its Uniform Resource Identifier (URI) is used to identify a particular resource, and the set of user selected input. This data is used as input to a python class called *Resource*. Upon creating the Resource object, the user defined request is validated. In case there is an error in the data, for example a past date, or an invalid airport/city, the proccess does not continue, and Django responds with an error. On the other hand, if the validation is succesfull, the Resource object creates a second class object, called *Request* object. There are 3 types of request objects: the *single*, *round* and *multicity* trip. These objects are not to be confused with the respective single/round flight. While a flight corresponds to a single instance connecting two locations in a particular date, a single/round flight may have an extended start period, in which case there are several flights which may constitute a solution to a user request.

Every Request class has a particular set of functions which are executed each time the class is instantiated. This set of instructions are called the *main* cycle of each request object, and correspond to the:

1. creation of a list of necessary flights;
2. acquisition of the necessary flight data;
3. construction of the weight matrix according to the objective function;
4. execution of the optimization algorithms;
5. construction of the solution;

From the above defined main cycle, steps i), iii) and v) are done internally, by each of the class object, while steps ii) and iv) are managed by the Data Management and Optimizaion system, respectively. Note that the step iv), the execution of optimization algorithms, is not necessary for the single/round trip objects, because the best solution is unambiguous.

Figure **??** illustrates the structure of the Client Side Application. This application is activated each time a request is received, which activates the instantiation of a *Resource* object, responsible for the validation of the request. After a successfull validation, the Resource object instantiates a *Request* object, which executes the *main cycle* of the solution construction procedure, by first calling the Data Management System to collect the necessary set of flights, following this with the execution of a series of optimization algorithms, to produce a stream of responses which can be served to respond to the request.
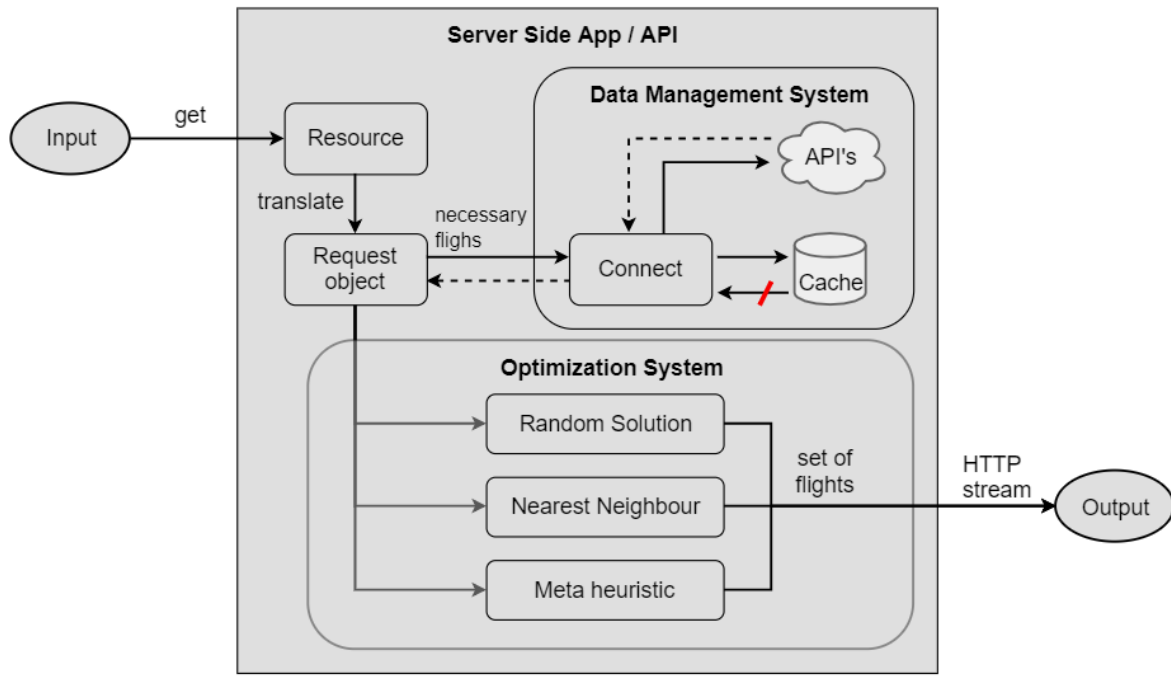
Figure 5.6: Structure of the Server Side Application/API.

5.3.2 Api protocol

This subsection introduces the syntax of the implemented API protocol. The objective of this protocol is to be simple and clear, and enable an easy interpretation of the user request.

Every user request may be formullated as a Flying Tourist Problem instance, as proposed in **??**. Thus, each request is characterized by a limited number of attributes (origin, start date, etc.). This means that the identification of a specific resource can be achieved by including each of these attributes in the uniform resource identifier.

Each request starts with the API endpoint, the specification of the resource type, '/flights', followed by the necessary atributes to describe the requeest. These atributes are grouped using a '&' symbol, and do not require a specific order. Table **??** specifies the possible request attributes, together with the keyword necessary to identify it, and the details and datatype of each.

5.3.3 Data Management System

The Data Management System (DMS) is responsible for collecting the set of necessary flights, in order to process user defined request. As discussed in section **??**, it is possible to obtain this flight data in two distint ways, by using a third-party flight data API, or by performing web scraping. Both of these methods were implemented and tested, and the results will be discussed in section **put the section of**

Table 5.2: This table specified the (keywords, value) pairs, which must be specified in order to uniquely identify a resource.

Name	Symbol	Keyword	Details
start city	v_0	flyFrom=	Requires a city name, or an ICAO code.
return city	v_{n+1}	returnTo=	Requires a city name, or an ICAO code.
destinations	V	cities=	Defines the cities to be visited. Accepts multiple values, sepparated by comma. Each city is specified by a city name or an ICAO code.
durations	D	duration=	Defines the duration of stay, in days, for each city. Each value must be a positive integer. Must be the same length as the number of cities.
start date	T_0	minDate = maxDate =	minDate specifies the earliest start date T_{0i}, while maxDateidentifies the max T_{0f}. Each date follows the dd/mm/yyyy format.

the comparison of API vs webscraping here. During the development of this work, several API's were tested, and the one which is currently being used is the *Kiwi.com* API, whose details can be found in the followig website: *https://docs.kiwi.com/*.

Communicating with a third-party API to request flight data is simply a matter of making HTTP requests using an URL which defines the resource under query. This request is usually answered with a JSON or XML object, which contains the relevant response for the performed request. In general, each API has its own URL syntax and response structure. Thus, communicating with different API's requires the differentiation of the resource identification and response parsing methods, because these are usually API dependent. Thus, communicating with an API usually involves three steps:

1. creation of a URL specifying the intended resource;
2. execution of an HTTP request to the URL;
3. deconstruction (parsing) of the response;

These three essential steps are the base for any data collection system. They can be used to collect data from an API, and they can also be used to do webscraping. In general, the difference between these two methods (api vs webscraping) is more likely to be felt in the parsing of the response. Using an API, the response is usually structured and organized, encoded in a JSON data type, or similar, which is, in general, human readable. In its turn, using web scraping, the response comes in the form of HTML, and the necessary data may be trapped under many levels of HTML objects. In general, it is harder to parse the result of webscraping.

There is a second difference that must be mentioned when comparing the usage of API's and webscraping. Webscraping is the act of retrieving the data visible on the *screen*. The problem is that, in many cases, before the data becomes visable, javascript has to be executed, and in some cases, there

are database acceses, which usually require a substaintial ammoun of time to execute. In this case, it is said that the websites uses javascript to render their page, and in general, webscraping javascript websites is much more harder and slower, because a simple HTTP request is not sufficient to obtain the necessary data. Instead, it is necessary to emulate a browser, make the HTTP request, and wait for the javascript to be completly loaded. During the development of this system, the *httplib* module of the python programming language was utilized for executing the HTTP requests, and the *Selenium Web driver* library for the execution of a automatable browser.

Given a list of flights whose data must be collected, figure **??** illustrates the necessary steps to communicate with a thid-party API, using HTTP protocol, to obtain the necessary data. In this figure, the system utilizes a serial approach, which means that at any time, only one request is being executed.
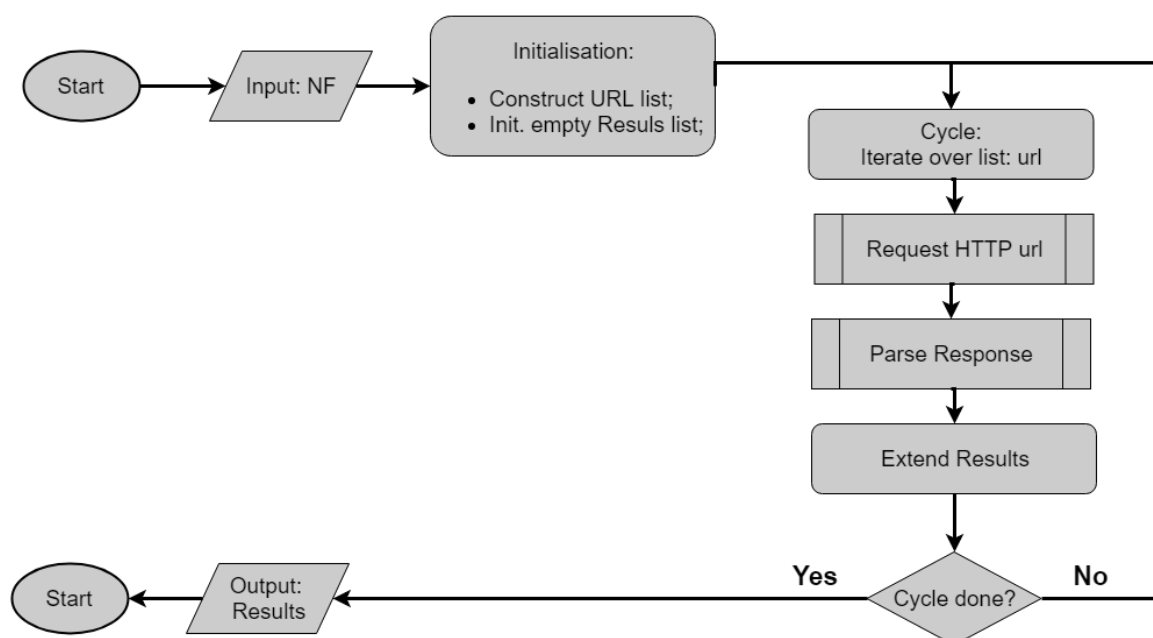
Figure 5.7: Communication with third-party API's, using HTTP protocol, and a serial requesting scheme.

The bottleneck of the serial system illustrated in figure **??**, is the necessary time to receive the response to an HTTP request. In order to take advantage of this bottleneck, a concurrent approach was considered, in which the waiting period of a request is utilized to spwan more requests. This approach is achieved by adopting a *Producer-Consumer* system, illustrated in figure **??**. This system spwans at most n_{max} threads (workers), one at a time, to execute a list of jobs, which correspond to making an HTTP request and parsing the response. Using this approach, the bottleneck experienced by each worker is not imposed on any of the other $n_{max} - 1$ workers, and thus the time-delay is not cummulative.

The webscraping systems developed are very similar to those illustrated in figures **??** and **??**. However, instead of making simple HTTP requests, a Selenium browser is used to access the website.
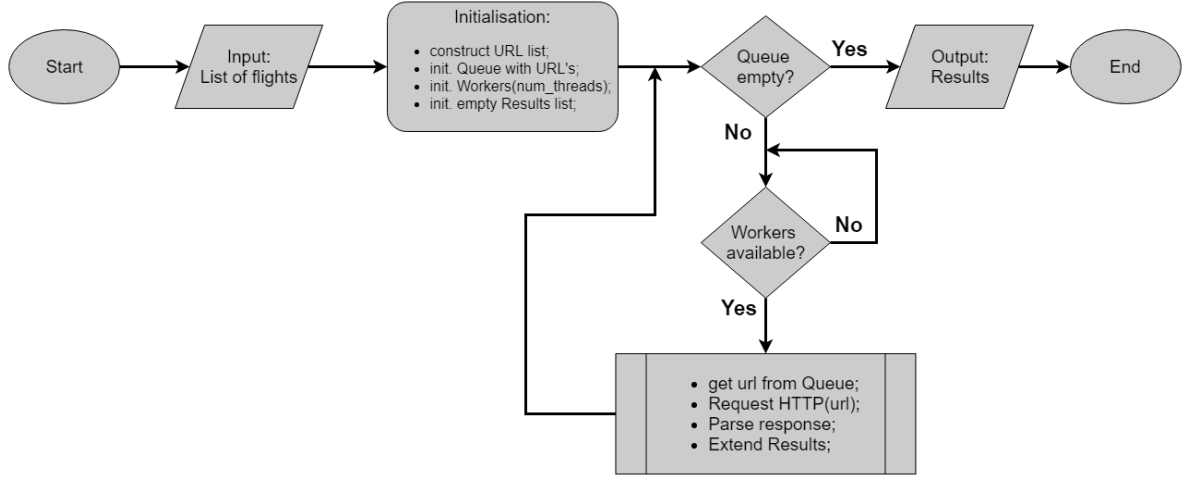
Figure 5.8: Communication with third-party API's, using HTTP protocol, and concurrent requesting scheme to take advantage of the waiting times.

Consequentially, the execution of the concurrent approach must create, at most, n_{max} browsers. The illustration of these systems may be consulted in the **ANEXO**.

Subsection **section herereeere** will evaluate the proposed system, and compare the efficiency of the serial and concurrent approaches, for both the API collectiong and the webscraping. These two systems will also be compared in order to evaluate which is most adequate for the usage in a production system.

5.4 Optimization System

Following the architecture proposed in section **??**, the optimization system will produce a stream of responses to the user request. To produce a fast response, a simple random solution will be generated initially. This is followed with the application of the nearest neighbour heuristic, and concludes with the execution of two meta-heuristic algorithms: the Ant Colony Optimization and Simmulated Annealing.

5.4.1 Initial Response

Consider a user request which defines a start period $T_0 = [T_{0i}, T_{0f}]$, a start city, v_0, a return city v_{n+1}, and a list of cities to be visited, V, where upon visiting each city $i \in V$, a waiting periof of d_i days is necessary.

The steps necessary to construct a random solution to this request are illustrated in figure **??**, and can be summarized as follows. Start by setting the initial time, $t \in T_0$, and the current node, $v_c = v_0$. If there are no nodes to visit, that is, if $V = \{\}$, a solution if defined by the flight $a^t_{v_0, v_{n+1}}$. Otherwise,

select a random city, $v_i \in V$, and extend the solution with the flight $a^t_{v_c, v_i}$. Follow this with the update of $t = t + d_i$, remove the selected node v_i from V, and set it as the current node. Repeat this process untill all nodes are visited, and conclude by closing the tour.

The described process may start immediatly after receiving a request, because it does not require the information relative to the arcs of the problem. Instead, it generates a completly random solution, composed of several flights, whose information is not yet known. Thus, this solution is passed to the Data Management System, which calls a third-party API to request the relevant information. If the number of flights which constitute the solution is lower than 9, this information can be obtained with a single HTTP request to the API. Otherwise, the number of necessary requests is given by, approximately, $N/10 + 1$, where N is the number of necessary flights.

Figure 5.9: Steps to generate a random solution to a user request.

While the random solution generation may start immediatly after receiving a request, the nearest neighbour heuristic relies on a complete cost matrix, and so can not be initiated before collecting all the information necessary. The nearest neighbour algorithm is closely related to the process described in figure **??**, however, instead of selecting a random node $i \in V$, this node is selected according to its cost. At any point of the process, there is a *current* node and time, which can be used to determine all the possible solution components (arcs), where the one with the lowest cost is selected. Furthermore, if the start time of the request is given by a time window, the illustrated process can be repeated for all allowable start dates.

5.4.2   Ant Colony Optimization

Figure **??** presents a block diagram illustrating the Ant Colony Optimization metaheuristic. The algorithm receives a weight matrix which contains the information relative to the arcs of the bipartite graph

describing the problem, aswell as other relevant data, as the duration set D, associated to each city i, of the set of nodes to be visited, V. After the initialisation, and untill the termination condition is met, the algorithm performs a cycle, in which every agent constructs a solution to the problem, followed by a pherormone matrix update, to reflect the search experience of each ant.
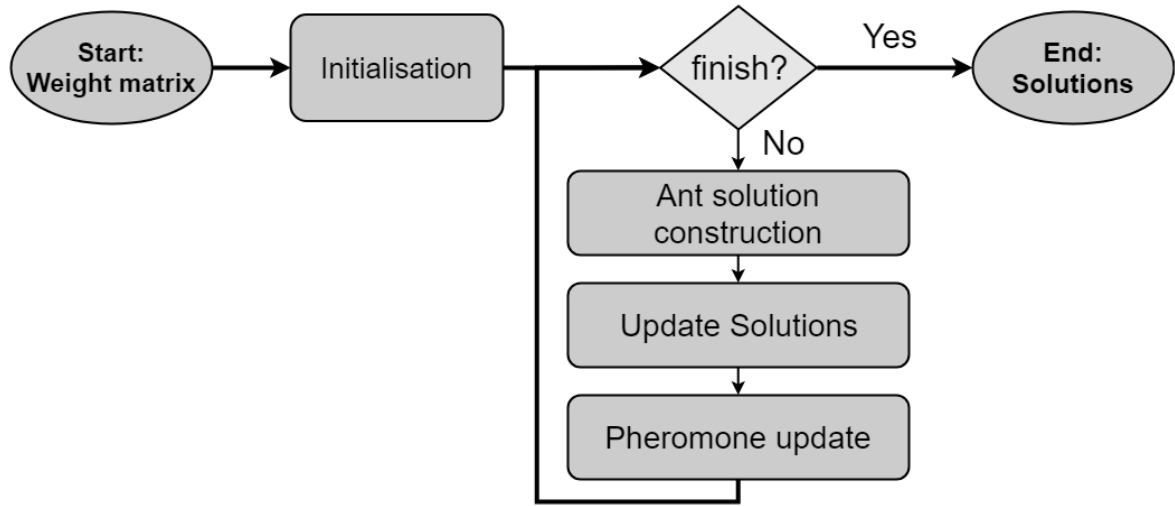

Figure 5.10: Block diagram of the Any Colony Optimization metaheuristic.

The initialisation of the ACO metaheuristic requires the construction of an initial pherormone matrix. As suggested by the ACO authors [**?**], the value of each entry of the pherormone matrix is given by equation **??**, which is inversely proportional to the cost of the nearest neighbour solution, C^{nn}. The initialisation of the metaheuristic also requires the definition of a variety of algorithm specific parameters, as the number of ants m, the pherormone evaporation rate ρ, the heuristic relative influence β, the pherormone relative influence α, and the exploration rate Q_0.

$$\tau_{ij}^t = \tau_0 = \frac{1}{nC^{nn}} \tag{5.1}$$

The construction process undertaken by each ant, illustrated in figure **??** is as follows. First, the current time is set to a value belonging to the allowable start dates, $t \in T_0$, and the current node is set to the start node v_0. Each ant enters than an iterative cycle untill all nodes belonging to V are visited. At every step of this cycle, an ant chooses a solution component by either *exploiting* or *exploring* the search space. The decision of exploiting or exploring depends on the algorithm parameter Q_0, and a pseudo-random value q, calculated at run time. The selection of the solution component j, which identifies the next city to be visited, is thus given by equation **??**. After the selection of each solution component, it is necessary to update the time, incrementing it by the duration relative to the selected city.

Figure 5.11: Block diagram of the construction procedure undertaken by each ant.

$$j = \begin{cases} exploitation \text{ (eq. ??)}, & \text{if } q \leq Q_0 \\ exploration \text{ (eq. ??)}, & \text{otherwise} \end{cases} \tag{5.2}$$

The exploration of the search space utilizes the so called *random-proportional rule*, given by equation **??**, which determines the next solution component of the ants solution. Note that $J_k(i,t)$ is the set of the solutions components which might be selected and form a valid solution, by an ant which is currently at city i at time t. In its turn, the exploration is given by equation **??**, and $p_a(i,j,t)$ represents the probability of ant a, which is currently at the node i at time t, selecting j as the next node to visit. In the presented equations, η is the inverse of the weight matrix value.

$$argmax_{j \in J_k(i,t)}[\tau(i,j,t)][\eta(i,j,t)]^\beta \tag{5.3}$$

$$p_a(i,j,t) = \begin{cases} \dfrac{[\tau(i,j,t)][\eta(i,j,t)]^\beta}{\sum_{u \in J_k(i,t)}[\tau(i,u,t)][\eta(i,u,t)]^\beta}, & \text{if } j \in J_k(i,t) \\ 0, & \text{otherwise} \end{cases} \tag{5.4}$$

Follwing the iterative construction procedure, an uncomplete solution is available. To complete this solution, it is necessary to add an extra solution component, which connects the last visited node, to the return node, v_{n+1}.

Having a complete solution, it is possible to perform *Daemon Actions*, or in other words, try to improve

the constructed solution using a local search algorithm. The execution of this step is optional, and is set by a parameter algorithm at the initialization of the procedure. The selection of a local search algorithm is usually problem specific, and for the Traveling Salesman Problem, one of the most used local search procedures is the k-opt exchange. During the development of this work, the 2-opt procedure was utilized.

The mainstream success of the Ant Colony Optimization in the multiple areas in which it was applied, is due to its ability to *guide* the agents to good solutions. While the construction of each solution follows a pseudo-random rule, the overall quality of this solution is used to bias the search towards the search space around this solution. This achieved by the global pheromone update, given by equation **??**.

$$\tau_{ijt} = (1 - \rho)\tau_{ijt} + \sum_{s \in S_{upd}|c_{ijt} \in s} g(s) \tag{5.5}$$
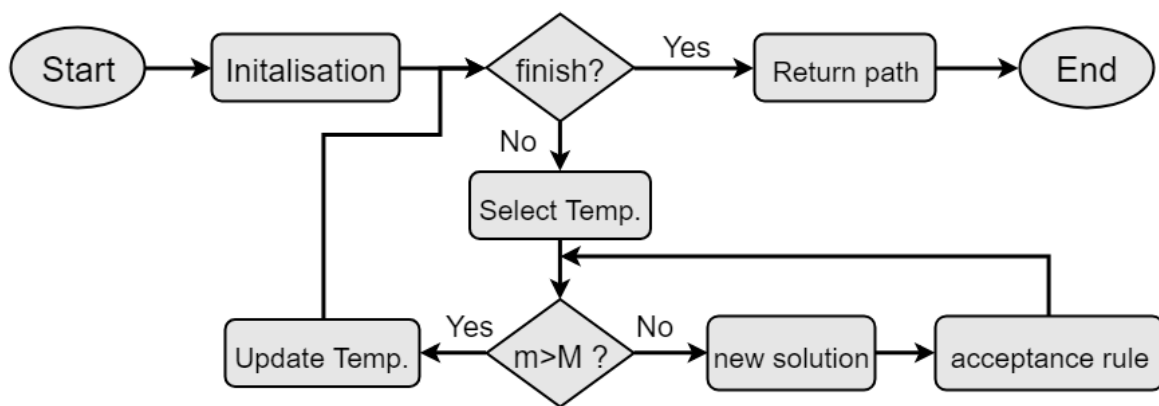
The global pheromone update is a proccess in which, at every step of the algorithm, all entries of the pheromone matrix are updated, by the so called pheromone *evaporation* and *depositation*. Pheromone evaporation is a method used to avoid getting stuck in sub-optimal solutions, while pheromone deposition intends to induce the exploration of the search space around good solutions, by incrementing the pheromone values of the solution components belonging to a set S_{upd}. The set S_{upd} is usually a subset of the solutions constructed at the current iterative step, $S_{upd} \subset S_{iter}$. Furthermore, the set S_{upd} includes the best global solution S_{best}, and this work follows the elitist ant rule, which defines that pheromone deposition on the solution components beloning to S_{best} is actually higher than those of the other solutions in S_{upd}. In equation **??**, $g() : S \rightarrow \mathcal{R}^+$ is a function such that $f(s) < f(s^{'}) \Rightarrow g(s) \geq g(s^{'})$. It is common set $g()$ to be inversely proportional to the objective function $f()$. Thus, the inversely proportional constant for g is set to 1 for all solution components in S_{upd}, and to a higher value for S_{best}, 3 in this case, due to the elitist ant rule.

5.4.3 Simmulated Annealing

The Simulated Annealing is a metaheuristic for solving continuous and discrete combinatorial optimization problem, and it is well known for its ability to escape from local minima. Every SA algorithm is composed of a main cycle, and a secondary cycle, which is often refered to as the Markov cycle. During the markov cycle, new solutions are generated, based on the current solution, and, if the new solution is better than the current one, it is accepted, while a worse solution is conditionally accepted, according to the Metropolis acceptance criterion. This key characteristic of accepting worse solutions, also called hill-climbing moves, is what enables the metaheuristic with the possibility of escaping local minima. The probability of accepting a worse solution is influenced by a parameter of the algorithm, called temperature, and is higher for higher values of the temperature. Thus, the metaheuristic is defined in such a way that the initial temperature is high, accepting worse solutions with a high probability, but during the

execution of the algorithm the temperature is decreased, and so is the probability of accepting a worse solution. The adjustmement of the temperature is done at each iteration of the main cycle, after running a complete markov cycle. As the algorithm reaches the end condition, the algorithm becomes gradually more gready, and accepts only improvements to the current solution.

Figure **??** presents a simplified block diagram for the Simulated Annealing metaheuristic. This block diagram is valid for the majority of the SA algorithms, because it is constituted by only the main blocks of the algorithm, and does not specify the cooling schedule, neither the solution generation process, nor the acceptance criteria. This three blocks of the diagram are responsible for introducing differentiation between the different SA algorithms.

Figure 5.12: Block diagram of the Simmulated Annealing metaheuristic.

The majority of the SA algorithms operate on a single solution, but there are authors which propose a multi-agent approach [**?**], and mention that the classic SA does not learn from its search history in an inteligent way, as other meta-heuristics, like the ACO, does. Other authors focus on creating a cooling schedule which may benefit the SA algorithm, by enabling a more exhaustive exploration of the search space around more promosing temperatures. An example of this is the List-Based SA algorithm proposed in [**?**], which creates a cooling schedule which initially decreases faster than the traditional geometric schedule, escaping faster from the non promising temperatures, and which than decreases slowly, inducing a more exhaustive search around promising temperatures. The implementation of the Simmulated Annealing algorithm developed in this work, follows the multi-agent and list-based chooling schedule, introduced in [**?**] and [**?**].

Following the block diagram illustrated in fig **??**, the algorithm receives a weight matrix describing the problem, and other relevant information, as the duration associated to each city, and possible constraints relating the initial and final node. The algorithm is than initialised, and some parameters must be set. This includes the initial acceptance probability p_0, the maximum length of the temperature list L_{max}, the main cycle stop criteria and the markov cycle length. The initialisation also requires the construction of

an initial temperature list, which is done as follows:

1. Create an empty temperature list L;
2. Create an initial solution x;
3. Create a candidate solution y from x;
4. If $f(y) < f(x)$, swap x and y;
5. Insert $t = \frac{-|f(y)-f(x)|}{p_0}$ into the temperature list;
6. Repeat steps iii) to v) untill the temperature list reaches its maximum length;

The Simmulated Annealing metaheuristic is based on two cycles. The inner cycle, often called the *Markov chain*, is responsible for producing candidate solutions based on the current one. The algorithm may accept or reject this candidate solution, and this depends on the solutions objective function value, aswell as the current state temperature. The Markov cycle usually runs for a fixed number of times, at each main iteration cycle. After completing the Markov cycle, the temperature is decreased, and if the termination condition is not met, the Markov cycle restarts.

The process by which a candidate solution is generated is problem specific. For the Traveling Salesman Problem, there are several suggestions in the literature, which include, but are not limited to, 2 and 3-opt moves, insertion, reversion and swapping mechanisms. During the development of this work, two stategies were tested. The first follows a simple 2-opt move strategy at each markov cycle, while the second is a greedy approach which selects the best result of the insertion, reversion and swapping functions.

At each step of the markov chain, the algorithm dictates that if a candidade solution is better than the current one, this solution is accepted, and set as the current one. In its turn, if a candidate solution is worse, it not automatically discarded, but it may be accepted. This set of rules is refered to as the Metropolis acceptance criteria, and is defines in equation **??**. The criteria sets that when confronted with a worse solution, a random number $r : r \in [0, 1[$ is generated, and if r is less than the acceptance probability, $e^{-\frac{f(y)-f(x)}{t}}$, the candidate solution is accepted. This probability is such that, the lower the objective function difference, the higher the probability of accepting the solution. On the contrary, as the state temperature decreases, so does the acceptance probability.

$$p = \begin{cases} 1, & \text{if } f(y) \leq f(x), \\ e^{-\frac{f(y)-f(x)}{t}}, & \text{otherwise} \end{cases} \qquad (5.6)$$

The Simulated Annealing convergence theory suggests that this algorithm is capable of reaching the global minima,

6

Experimental Results

Contents

6.1 Experimental Results

In order to evaluate the performance of the developed system, several tests, with different goals, were developed and executed. First and foremost, the optimization algorithms are tested in order to evaluate their overall quality on a set of benchmark tests. This is followed by a set of tests over the Data Management System, in order to estimate the necessary time to collect the required data to answer to any given request. Finally, the overall utility of the proposed and developed system is evaluated, by performing a series of tests on the Flying Tourist Problem, and comparing the solutions provided by the metaheuristic to those from *random* optimization algorithm, and the nearest neighbor heuristic.

This experiments were executed on a 2.6GHz Intel i7-6700, and all the code was developed using the Python programming language.

6.1.1 Optimization system

Given that there are no benchmark tests available for the time-dependent TSP, we tested the algorithms on the classical Traveling salesman. This was followed by a set of tests on TDTSP instances, created by the method described in ([**?**]), which uses dual-theory on the Integer Linear formulation for the TDTSP, to generate TDTSP instances whose optimal solution is defined a-priori.

In order to test the performance of the optimization system, 5 set of problems were solved, ranging from 17 to 323 nodes, each problem was executed 5 times, and its results averaged. Note that each algorithm was allowed to run for a maximum of 60 seconds. The choice of executing the metaheuristic for a limited amount of time, instead of a fixed number of iterations, arises from the necessity to produce time-efficient responses to the user. The results of the execution of these tests are presented in table **??** and **??**, for the asymmetric TSP and time-dependent TSP, respectively.

The analysis of table **??**, regarding the resolution of TSP benchmark problems, gives some interesting information. For small instances (17 nodes), both ACO and SA present the optimal solution, consistently. For medium instances (35-53 nodes), both algorithm perform in the range of 5-20% relative error. As for bigger problems (170-323 nodes), the performance of the ACO decreases only slightly (22-25%), while that of the SA becomes much worse (29-63%).

The performance of the Simulated Annealing on the test problem $ftv170$, which appears to be somewhat abnormal, due to its very high error, was further analyzed. The tests were repeated 2 more times, averaging the result of 5 execution, for a total duration of 60 seconds each. The quality of the results did not improve. So, instead of running the algorithm for just 60 seconds, we executed a series of tests were it run for a longer time, 180 seconds. This time, the results were somewhat better, presenting a relative error of 45.72%, as opposed to the previous 63%. While this improvement is still low from the desirable amount, it suggest that, with enough time, the Simulated Annealing could possible present much better

Table 6.1: Performance of the metaheuristics on the **Asymmetric TSP**.

algorithm	problem	opt cost	meta cost	num. iterations	rel. error
SA	br17	39	39	264319.4	0
	ftv35	1473	1641	63659	11.41
	ft53	6905	7963	36189.2	15.32
	ftv170	2755	4486.4	6714.4	62.84
	rbg323	1326	1706.2	1361.8	28.67
ACO	br17	39	39	6892.7	0
	ftv35	1473	1552.4	1469.6	5.39
	ft53	6905	8269.6	689.2	19.76
	ftv170	2755	3359.4	67.8	21.93
	rbg323	1326	1660.4	19.6	25.21

solutions.

Looking now at the performance over the time-dependent TSP, presented in table **??**, the results are very different from thus of the classic TSP. First of all, no optimal solution could be consistently found. Furthermore, the relative error was only above 10% once, and this occurs for the smallest instance (p17), using the ACO algorithm. (This suggests that the algorithm is stuck, using to much heuristic information, and low exploration. Given that the heuristic parameter β, is very high, this is not surprising. The confirmation of this is given by the better performance on the bigger instances.)

The comparison of the SA vs the ACO deserves some extra notes. First of all, it is worth nothing that, at each iteration of both algorithms, the same number of solutions are created, since the markov chain length, and the number of ants, are both set to N, the number of nodes. A comparison of the number of iterations executed by each algorithm, over the 60 available seconds, leads to the conclusion that SA performs much more iterations than the ACO. This should come as no surprise, given that the SA relies on local search and improvement heuristic, while ACO requires the *construction* of solutions. This for itself is not the problem. The problem is that the solution construction process of the ACO is somewhat complex and computational heavy. At each ant construction step, several difficult calculations have to be performed, as to select the following solution component.

Despite being heavier and slower, the Ant Colony System appears to be, overall, slightly better than the Simulated Annealing, even without the use of local search procedures. Another advantage of the Ant Colony Optimization is that good solutions are produced very early, and not only in the late stages of the algorithm execution, as occurs with the Simulated Annealing. This means that an early interruption in the execution of the ACO may still produce good results, while in the Simulated Annealing this becomes very unlikely.

Table 6.2: Performance of the optimization algorithm on the **time-dependent TSP**.

| algorithm	problem	opt cost	meta cost	num. iterations	rel. error
SA	p17	2458	2631.2	219517.2	7.04
	p35	5131	5406.8	80262	5.38
	p53	7930	8265.6	37450.6	4.23
	p170	25483	26427	3521	3.70
	p323	48991	51926.8	900.6	5.99
ACO	p17	2458	2729	9720.6	11.03
	p35	5131	5500.2	2590.6	7.20
	p53	7930	8370.4	1099.6	5.53
	p170	25483	26402.4	71.2	3.61
	p323	48991	50261.4	13.4	2.59

6.1.2 Utility Test

In order to establish the actual utility of the proposed system, the following experiment was executed. A series of Flying Tourist Problem (FTP) instances were constructed, ranging from just 1 city to visit (which corresponds to a round flight - if the return node is equal to the start node), up to a total of 20 cities. For each problem instance, solutions were constructed using three different algorithms: the random algorithm (explained in the introduction of section **??**); the nearest neighbor heuristic; and both ACO and SA metaheuristics.

For every different instance size, 5 different FTP were considered, and the results averaged. In all cases, the trip starts and returns to Lisbon, and visits a given set of cities, chosen randomly, belonging to (Miami, Moscow, JFK, Hong Kong, Sidney, Dubai, Kiev, Barcelona, Madrid, Dublin, Los Angeles, Beijing, San Francisco, Singapore, Johannesburg, Cairo, Casablanca, Abuja, Frankfurt, Atlanta, Istanbul, Oslo). The start date was always the same, set to *1 July 2018*, which, upon the execution of the tests, was a date 45 days into the future. Finally, the waiting period on *each* city was set to a random value between 1 and 12 days.

The results provided by the different algorithms are presented in table **??**, and include a comparison of the flight cost and flight duration provided by the different algorithms. Note that R refers to the random algorithm, while NN to the nearest neighbor, and MC and ME to the metaheuristics, utilizing the cost (eq. **??**) and entropy (eq. **??**) objective functions, respectively. This table also presents the completion time of each of these algorithms.

Note that the nearest neighbor and metaheuristics run only after the complete weight matrix is available, which occurs only once the communication with the third-party API is finished. Note also that, after the weight matrix is complete, the nearest neighbor is run first, followed by the metaheuristic for the cost, and finally by the metaheuristic for the entropy. Thus, the completion time of the nearest neighbor, gives a good estimation, or "upper bound", on the time spent communicating with the third-party API, as to obtain the necessary flight data. Figure **??** presents an accurate illustration of the time spent

86

Figure 6.1: Time spent communication with the Kiwi API, as to obtain the necessary flight data, as a function of the number of nodes, and length of the start-period.

Figure 6.2: Comparison of the solution cost, of the nearest neighbor and metaheuristic, relative to the initial solution produced, for a single start date.

communicating with the third-party API, as a function of the number of nodes, and the length of the start-period.

As to simplify the comprehension of the results presented in table **??**, figure **??** illustrates the improvement of the total flight cost, for the NN and metaheuristic algorithms, relative to the cost of the initial *random* solution. It is worth noting that, for a single node (round flight), the initial random solution, the nearest neighbor, and the metaheuristic, all produce the same result. For 3 nodes, the nearest neighbor does not present relevant improvement, while the metaheuristic already distances itself, presenting around 5% improvement. For 5 nodes, this difference increases to 20%, and at 10 nodes it reaches the 50% mark. Instances with more nodes, 15-20, continue to improve, slowly, up to 60%.

We also analyzed the impact of using different objective functions in the optimization process. Instead of optimizing for the total trip cost, the algorithm was set to minimize the entropy, that is, the weighted summation of the flight cost and flight duration. The entropy is set such that the cost weight contributes with 70%, and the flight duration with the remaining 30%. After executing the metaheuristic for both objective functions, the total flight cost and duration are compared. The results are illustrated in figure

Figure 6.3: Comparison of the variation of the total flight cost and duration, for two different objective functions, minimizing the total cost and entropy.

Figure 6.4: Comparison of the solution cost, of the nearest neighbor and metaheuristic, relative to the initial solution produced, for a start-period of 31 days.

??. The inspection of this figure leads to the conclusion that, in general, by increasing the cost by around 20%, the flight duration decreases to approximately half.

This series of experiments is completed with the execution of the same problems, this time with a start-period of 31 days. The impact in the solution improvement, relative to the trip cost, can be verified in figure **??**. The analysis of this figure leads to the conclusion that, the selection of an extended start-period, as opposed to a single date, leads to a higher increase in the total savings. For the round flight, the improvement of the metaheuristic is around 10%. This is because the random and the nearest neighbor solution consider a single start date for the solution construction process. Other small instances, with 3 and 5 nodes, present much higher improvements than those with a single start date, ranging from 20-40%, as opposed to just 5-20%. For bigger problem instances, the effect of increasing the start-period is less significant, but it exists, providing an improvement of up to 65%.

Table 6.3: Comparison of the flight cost, flight duration, and execution time, for Flying Tourist Problem instances, ranging from 1 to 20 nodes, for a single start date, using three different algorithms.

—N—	Cost (Euro)				Fly duration (hours)				Completion time (seconds)			
	R	NN	MC	ME	R	NN	MC	ME	R	NN	MC	ME
1	838.2	838.2	838.2	1158.4	82.11	82.11	82.11	37.37	2.57	5.03	7.87	12.13
3	1520.6	1512.8	1457.8	1828.4	124.91	117.65	111.3	58.49	2.64	7.24	11.79	17.39
5	2088.6	1993	1654.8	2123.7	194.00	157.14	157.14	63.39	3.17	10.05	15.00	21.20
7	3571.2	2479	1927.6	2520.6	267.53	203.16	180.59	90.35	3.43	13.48	18.48	25.16
10	4819	3228.8	2468.2	3211.8	377.95	255.89	236.82	105.06	3.69	25.72	30.73	38.31
15	7288.8	3808.2	3103.8	3996.6	518.21	312.66	259.83	141.81	4.59	55.33	60.37	68.91
20	9187.6	4664.6	3721.4	4704.2	665.85	343.57	324.62	162.71	5.38	133.62	138.66	148.67

6.2 Conclusions

Placeholder for the conclusions.

A

Code of Project

Nulla dui purus, eleifend vel, consequat non, dictum porta, nulla. Duis ante mi, laoreet ut, commodo eleifend, cursus nec, lorem. Aenean eu est. Etiam imperdiet turpis. Praesent nec augue. Curabitur ligula quam, rutrum id, tempor sed, consequat ac, dui. Vestibulum accumsan eros nec magna. Vestibulum vitae dui. Vestibulum nec ligula et lorem consequat ullamcorper.

Listing A.1: Example of a XML file.

```
1  <?xml version="1.0" encoding="UTF-8"?>
2  <StreamInfo version="2.0">
3      <Clip duration="PT01M0.00S">
4          <BaseURL>videos/</BaseURL>
5          <Description>svc_1</Description>
6          <Representation mimeType="video/SVC" codecs="svc" frameRate="30.00" bandwidth="401.90"
7              width="176" height="144" id="L0">
8              <BaseURL>svc_1/</BaseURL>
9              <SegmentInfo from="0" to="11" duration="PT5.00S">
```

```xml
10          <BaseURL>svc_1-L0-</BaseURL>
11        </SegmentInfo>
12      </Representation>
13    <Representation mimeType="video/SVC" codecs="svc" frameRate="30.00" bandwidth="1322.60"
14        width="352" height="288" id="L1">
15      <BaseURL>svc_1/</BaseURL>
16      <SegmentInfo from="0" to="11" duration="PT5.00S">
17        <BaseURL>svc_1-L1-</BaseURL>
18      </SegmentInfo>
19    </Representation>
20  </Clip>
21 </StreamInfo>
```

Etiam imperdiet turpis. Praesent nec augue. Curabitur ligula quam, rutrum id, tempor sed, consequat ac, dui. Maecenas tincidunt velit quis orci. Sed in dui. Nullam ut mauris eu mi mollis luctus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Sed cursus cursus velit. Sed a massa. Duis dignissim euismod quam.

Listing A.2: Assembler Main Code.

```asm
1  ; *************************************************************************
2  ; * Constantes
3  ; *************************************************************************
4
5  ON      EQU 1 ; contagem ligada
6  OFF     EQU 0 ; contagem desligada
7  INPUT   EQU 8000H ; endereço do porto de entrada
8          ;(bit 0 = RTC; bit 1 = botão)
9  OUTPUT  EQU 8000H ; endereço do porto de saída.
10
11
12 ; *************************************************************************
13 ; * Stack
14 ; *************************************************************************
15
16 PLACE    1000H
17 pilha:    TABLE 100H  ; espaço reservado para a pilha
18 fim_pilha:
19
20 ; *************************************************************************
21
22 PLACE    2000H
23
24 ; Tabela de vectores de interrupção
25
26 tab:    WORD   rot0
27
28 ; *************************************************************************
29 ; * Programa Principal
30 ; *************************************************************************
31
32 PLACE    0
33
34 inicio:
35   MOV BTE, tab     ; incializa BTE
36   MOV R9, INPUT    ; endereço do porto de entrada
37   MOV R10, OUTPUT    ; endereço do porto de Ìsada
38   MOV SP, fim_pilha
39   MOV R5, 1     ; inicializa estado do processo P1
40   MOV R6, 1     ; inicializa estado do processo P2
41   MOV R4, OFF     ; inicializa controle de RTC
42   MOV R8, 0     ; inicializa contador
43   MOV R7, OFF     ; inicialmente não permite contagem
44   EIO           ; permite interrupções tipo 0
```

94

```asm
45   EI            ; activa interrupções
46
47 ciclo:
48   CALL  P1      ; invoca processo P1
49   CALL  P2      ; invoca processo P2
50   JMP   ciclo   ; repete ciclo
51
52 ; ***********************************************************************
53 ;* ROTINAS
54 ; ***********************************************************************
55
56 P1:
57   CMP R5, 1    ; se estado = 1
58   JZ  P1_1
59   CMP R5, 2    ; se estado = 2
60   JZ  P1_2
61 sai_P1:
62   RET          ; sai do processo.
63
64
65 P1_1:
66   MOVB  R0, [R9]  ; lê porto de entrada
67   BIT R0, 1
68   JZ  sai_P1     ; se botão não carregado, sai do processo
69   MOV R7, ON     ; permite contagem do display
70   MOV R5, 2      ; passa ao estado 2 do P1
71   JMP sai_P1
72
73 P1_2:
74   MOVB  R0, [R9]  ; lê porto de entrada
75   BIT R0, 1
76   JNZ sai_P1     ; se botão continua carregado, sai do processo
77   MOV R7, OFF    ; caso contrário, desliga contagem do display
78   MOV R5, 1      ; passa ao estado 1 do P1
79   JMP sai_P1
```

Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Phasellus eget nisl ut elit porta ullamcorper. Maecenas tincidunt velit quis orci. Sed in dui. Nullam ut mauris eu mi mollis luctus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos.

This inline MATLAB code `for i=1:3, disp('cool'); end;` uses the \mcode{} command.[1]

Nullam ut mauris eu mi mollis luctus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Sed cursus cursus velit. Sed a massa. Duis dignissim euismod quam. Nullam euismod metus ut orci.

Listing A.3: Matlab Function

```matlab
1 for i = 1:3
2     if i >= 5 && a ~= b       % literate programming replacement
3         disp('cool');         % comment with some LaTeXin it:  πx²
4     end
5     [:,ind] = max(vec);
6     x_last = x(1,end) - 1;
7     v(end);
8     ylabel('Voltage (μV)');
9 end
```

---

[1] MATLAB Works also in footnotes: `for i=1:3, disp('cool'); end;`

Nullam ut mauris eu mi mollis luctus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Sed cursus cursus velit. Sed a massa. Duis dignissim euismod quam. Nullam euismod metus ut orci.

Listing A.4: `function.m`

```matlab
1  %    Copyright 2010 The MathWorks, Inc.
2  function ObjTrack(position)
3  % #codegen
4  % First, setup the figure
5  numPts = 300;                % Process and plot 300 samples
6  figure;hold;grid;            % Prepare plot window
7  % Main loop
8  for idx = 1: numPts
9      z = position(:,idx);     % Get the input data
10     y = kalmanfilter(z);     % Call Kalman filter to estimate the position
11     plot_trajectory(z,y);    % Plot the results
12 end
13 hold;
14 end   % of the function
```

Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Phasellus eget nisl ut elit porta ullamcorper. Maecenas tincidunt velit quis orci. Sed in dui. Nullam ut mauris eu mi mollis luctus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Sed cursus cursus velit. Sed a massa. Duis dignissim euismod quam. Nullam euismod metus ut orci. Vestibulum erat libero, scelerisque et, porttitor et, varius a, leo.

Listing A.5: HTML with CSS Code

```html
1  <!DOCTYPE html>
2  <html>
3    <head>
4      <title>Listings Style Test</title>
5      <meta charset="UTF-8">
6      <style>
7        /* CSS Test */
8        * {
9          padding: 0;
10         border: 0;
```

```
11          margin: 0;
12        }
13      </style>
14      <link rel="stylesheet" href="css/style.css" />
15  </head>
16  <header> hey </header>
17  <article> this is a article </article>
18  <body>
19    <!-- Paragraphs are fine -->
20    <div id="box">
21      <p>
22        Hello World
23      </p>
24      <p>Hello World</p>
25      <p id="test">Hello World</p>
26      <p></p>
27    </div>
28    <div>Test</div>
29    <!-- HTML script is not consistent -->
30    <script src="js/benchmark.js"></script>
31    <script>
32      function createSquare(x, y) {
33        // This is a comment.
34        var square = document.createElement('div');
35        square.style.width = square.style.height = '50px';
36        square.style.backgroundColor = 'blue';
37
38        /*
39         * This is another comment.
40         */
41        square.style.position = 'absolute';
42        square.style.left = x + 'px';
43        square.style.top = y + 'px';
44
45        var body = document.getElementsByTagName('body')[0];
46        body.appendChild(square);
47      };
48
```

97

```
49        // Please take a look at +=
50        window.addEventListener('mousedown', function(event) {
51          // German umlaut test: Berührungspunkt ermitteln
52          var x = event.touches[0].pageX;
53          var y = event.touches[0].pageY;
54          var lookAtThis += 1;
55        });
56      </script>
57    </body>
58  </html>
```

Nulla dui purus, eleifend vel, consequat non, dictum porta, nulla. Duis ante mi, laoreet ut, commodo eleifend, cursus nec, lorem. Aenean eu est. Etiam imperdiet turpis. Praesent nec augue. Curabitur ligula quam, rutrum id, tempor sed, consequat ac, dui. Vestibulum accumsan eros nec magna. Vestibulum vitae dui. Vestibulum nec ligula et lorem consequat ullamcorper.

Listing A.6: HTML CSS Javascript Code

```
1
2  @media only screen and (min-width: 768px) and (max-width: 991px) {
3
4    #main {
5      width: 712px;
6      padding: 100px 28px 120px;
7    }
8
9    /* .mono {
10     font-size: 90%;
11   } */
12
13   .cssbtn a {
14     margin-top: 10px;
15     margin-bottom: 10px;
16     width: 60px;
17     height: 60px;
18     font-size: 28px;
19     line-height: 62px;
20   }
```

Nulla dui purus, eleifend vel, consequat non, dictum porta, nulla. Duis ante mi, laoreet ut, commodo eleifend, cursus nec, lorem. Aenean eu est. Etiam imperdiet turpis. Praesent nec augue. Curabitur ligula quam, rutrum id, tempor sed, consequat ac, dui. Vestibulum accumsan eros nec magna. Vestibulum vitae dui. Vestibulum nec ligula et lorem consequat ullamcorper.

Listing A.7: PYTHON Code

```python
class TelgramRequestHandler(object):
    def handle(self):
        addr = self.client_address[0]          # Client IP-adress
        telgram = self.request.recv(1024)      # Recieve telgram
        print "From: %s, Received: %s" % (addr, telgram)
        return
```

B
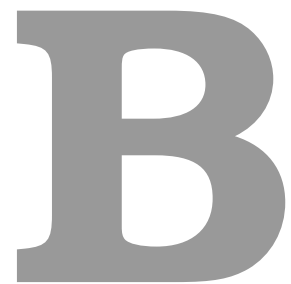
A Large Table

Aliquam et nisl vel ligula consectetuer suscipit. Morbi euismod enim eget neque. Donec sagittis massa. Vestibulum quis augue sit amet ipsum laoreet pretium. Nulla facilisi. Duis tincidunt, felis et luctus placerat, ipsum libero vestibulum sem, vitae elementum wisi ipsum a metus. Nulla a enim sed dui hendrerit lobortis. Donec lacinia vulputate magna. Vivamus suscipit lectus at quam. In lectus est, viverra a, ultricies ut, pulvinar vitae, tellus. Donec et lectus et sem rutrum sodales. Morbi cursus. Aliquam a odio. Sed tortor velit, convallis eget, porta interdum, convallis sed, tortor. Phasellus ac libero a lorem auctor mattis. Lorem ipsum dolor sit amet, consectetuer adipiscing elit.

Nunc auctor bibendum eros. Maecenas porta accumsan mauris. Etiam enim enim, elementum sed, bibendum quis, rhoncus non, metus. Fusce neque dolor, adipiscing sed, consectetuer et, lacinia sit amet, quam. Suspendisse wisi quam, consectetuer in, blandit sed, suscipit eu, eros. Etiam ligula enim, tempor ut, blandit nec, mollis eu, lectus. Nam cursus. Vivamus iaculis. Aenean risus purus, pharetra in, blandit quis, gravida a, turpis. Donec nisl. Aenean eget mi. Fusce mattis est id diam. Phasellus faucibus interdum sapien. Duis quis nunc. Sed enim. Nunc auctor bibendum eros. Maecenas porta accumsan mauris. Etiam enim enim, elementum sed, bibendum quis, rhoncus non, metus. Fusce neque dolor, adipiscing sed, consectetuer et, lacinia sit amet, quam.

Table B.1: Example table

Benchmark: ANN	#Layers (1)	#Nets (2)	#Nodes* $(3) = 8 \cdot (1) \cdot (2)$	Critical path $(4) = 4 \cdot (1)$	Latency (T_{iter}) (5)
A1	**3–1501**	1	**24–12008**	**12–6004**	4
A2	501	1	4008	2004	**2–2000**
A3	10	**2–1024**	**160–81920**	40	60†
A4	10	50	4000	40	**80–1200**

Benchmark: FFT	FFT size‡ (1)	#Inputs $(2) = 2^{(1)}$	#Nodes* $(3) = 10 \cdot (1) \cdot (2)$	Critical path $(4) = 4 \cdot (1)$	Latency (T_{iter}) (5)
F1	**1–10**	2–1024	**20–102400**	4–40	6–60†
F2	**5**	32	1600	20	**40 – 1500**

Benchmark: Random networks	#Types (1)	#Nodes (2)	#Networks (3)	Critical path (4)	Latency (T_{iter}) (5)
R1	3	10–2000	500	*variable*	(4)
R2	3	50	500	*variable*	$(4) \times [1; \cdots ; 20]$

* *Excluding constant nodes.*
† *Value kept proportional to the critical path:* $(5) = (4) * 1.5$.
‡ *A size of x corresponds to a 2^x point FFT.*
Values in bold indicate the parameter being varied.

As **??** shows, the data can be inserted from a file, in the case of a somehow complex structure. Notice the Table footnotes.

Lorem ipsum dolor sit amet, consectetuer adipiscing elit. Morbi commodo, ipsum sed pharetra gravida, orci magna rhoncus neque, id pulvinar odio lorem non turpis. Nullam sit amet enim. Suspendisse id velit vitae ligula volutpat condimentum. Aliquam erat volutpat. Sed quis velit. Nulla facilisi. Nulla libero. Vivamus pharetra posuere sapien. Nam consectetuer. Sed aliquam, nunc eget euismod ullamcorper, lectus nunc ullamcorper orci, fermentum bibendum enim nibh eget ipsum. Donec porttitor ligula eu dolor. Maecenas vitae nulla consequat libero cursus venenatis. Nam magna enim, accumsan eu, blandit sed, blandit a, eros.

And now an example (**??**) of a table that extends to more than one page. Notice the repetition of the Caption (with indication that is continued) and of the Header, as well as the continuation text at the bottom.

Table B.2: Example of a very long table spreading in several pages

	Time (s)	Triple chosen	Other feasible triples
0	(1, 11, 13725)	(1, 12, 10980), (1, 13, 8235), (2, 2, 0), (3, 1, 0)
2745	(1, 12, 10980)	(1, 13, 8235), (2, 2, 0), (2, 3, 0), (3, 1, 0)
5490	(1, 12, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
8235	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
10980	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)

Time (s)	Triple chosen	Other feasible triples
13725	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
16470	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
19215	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
21960	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
24705	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
27450	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
30195	(2, 2, 2745)	(2, 3, 0), (3, 1, 0)
32940	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
35685	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
38430	(1, 13, 10980)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
41175	(1, 12, 13725)	(1, 13, 10980), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
43920	(1, 13, 10980)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
46665	(2, 2, 2745)	(2, 3, 0), (3, 1, 0)
49410	(2, 2, 2745)	(2, 3, 0), (3, 1, 0)
52155	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
54900	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
57645	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
60390	(1, 12, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
63135	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
65880	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
68625	(2, 2, 2745)	(2, 3, 0), (3, 1, 0)
71370	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
74115	(1, 12, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
76860	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
79605	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
82350	(1, 12, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
85095	(1, 12, 13725)	(1, 13, 10980), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
87840	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
90585	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
93330	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
96075	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
98820	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
101565	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
104310	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
107055	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
109800	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
112545	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
115290	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
118035	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
120780	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
123525	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
126270	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
129015	(2, 2, 2745)	(2, 3, 0), (3, 1, 0)
131760	(2, 2, 2745)	(2, 3, 0), (3, 1, 0)
134505	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
137250	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
139995	(2, 2, 2745)	(2, 3, 0), (3, 1, 0)
142740	(2, 2, 2745)	(2, 3, 0), (3, 1, 0)
145485	(1, 12, 16470)	(1, 13, 13725), (2, 2, 2745), (2, 3, 0), (3, 1, 0)
148230	(2, 2, 2745)	(2, 3, 0), (3, 1, 0)

Table B.2 – continued from previous page

Time (s)	Triple chosen	Other feasible triples
150975	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
153720	(1, 12, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
156465	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
159210	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
161955	(1, 13, 16470)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)
164700	(1, 13, 13725)	(2, 2, 2745), (2, 3, 0), (3, 1, 0)